

AEA DISTINGUISHED LECTURE

Distorted Innovation: Does the Market Get the Direction of Technology Right?[†]

By DARON ACEMOGLU*

In the presence of markup differences and externalities, the equilibrium direction of innovation can be systematically distorted. I build a simple model of endogenous technology, which generalizes existing comparative statics and characterizes distortions in the direction of innovation. Empirical findings across a number of areas are consistent with this framework's predictions, and I use data from several studies to estimate its key parameters. Combining these numbers with estimates of externalities and markups, I find that equilibrium distortions in the direction of technology can be substantial and have large welfare costs in the context of industrial automation, health care, and energy.

There is broad agreement that technological change has been a major engine of economic growth and prosperity over the last 250 years.¹ However, not all innovations are created equal, and the *direction of technology* matters greatly as well.

Both antibiotics and dietary supplements have resulted from new innovations and have led to products that have been consumed by billions of people around the world. But most would agree

that antibiotics constitute a bigger technological breakthrough and have been socially more beneficial.² More strikingly, the same advances by early twentieth-century chemists, especially Fritz Haber and Carl Bosch, paved the way to both synthetic agricultural fertilizers, which massively boosted crop yields, and the large-scale production of more powerful explosives, which led to the death of millions of soldiers and civilians (e.g., Hager 2009). Few people would think that these two advances have similar social value. Additionally, different technologies often create gains and losses for different groups and may even influence other major social outcomes, including civic participation and democracy.

Economists have long recognized that the overall amount of research effort may be insufficient, and, as a result, government support for innovation—for instance, in the form of investments in the research infrastructure or

*Department of Economics, Massachusetts Institute of Technology (email: daron@mit.edu). This paper is based on the American Economic Association Distinguished Lecture delivered in New Orleans, January 2023. I thank Ufuk Akcigit, David Autor, Rena Conti, Joe Doyle, Michael Greenstone, Simon Johnson, Chad Jones, Amy Finkelstein, Will Rafey, Dani Rodrik, and John Van Reenen for very useful comments, discussions, and references. I am grateful to Juanita Jaramillo, Shinnosuke Kikuchi, Fredric Kong, and Todd Lensman for fantastic research assistance. Special thanks are due to David Hémous, Ralf Martin, Jacob Moscona, and John Van Reenen for sharing data and help with the empirical work reported in this paper. Last but not least, this work heavily draws on joint work with several coauthors. I am particularly grateful to Ufuk Akcigit, David Autor, Simon Johnson, Pascual Restrepo, and Fabrizio Zilibotti for their enduring contributions to my knowledge and understanding on these topics. All remaining errors are of course my own.

[†]Go to <https://doi.org/10.1257/pandp.20231000> to visit the article page for additional materials and author disclosure statement.

¹See, for example, Mokyr (1992, 2004) and Koyama and Rubin (2022).

²US annual expenditures are about \$10 billion in the 2010s for antibiotics and above \$30 billion for dietary supplements (see <https://academic.oup.com/cid/article/66/2/185/4093915> and <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3952619/>). This is despite the fact that there are no well-established studies documenting the effectiveness of dietary supplements (see <https://pubmed.ncbi.nlm.nih.gov/32601065/> and <https://www.nytimes.com/2016/11/15/well/eat/studies-show-little-benefit-in-supplements.html>).

R&D tax credits—is beneficial (see, e.g., Jones and Williams 1998; Bloom, Shankerman, and Van Reenen 2013; Howell 2017; and Azoulay et al. 2019 for evidence). Nevertheless, a common perspective is that the market is the best judge of how research efforts should be allocated: once basic support for research is provided, the government should have a limited role in influencing the direction of innovation. There are indeed myriad examples of failed government attempts at “picking winners” (see Pack and Saggi 2006 and Hufbauer and Jung 2021 for reviews). British science writer Ridley (2020) argues that the cumulative, step-by-step process of innovation is inevitably hampered when governments try to influence its direction. The opposite point of view emphasizes the myriad distortions in the equilibrium innovation process.

In this paper, I take an intermediate position. I assume that the market (working through competition between corporations and scientists) is best placed to experiment with new methods and carry out innovations, but the direction of technology can be systematically distorted. This position distinguishes my approach both from older-style industrial policy, where the government is assumed to have some ability to judge which sectors are more promising, and from more recent arguments claiming that the government can be as good as the private sector in innovation (e.g., Mazzucato 2015).³

To develop these ideas, I extend the directed technological change framework in Acemoglu (1998, 2002), focusing on an economy in which the private sector spearheads innovation and can target either one or both of two alternative,

imperfectly substitutable technologies. From a positive perspective, the framework links the direction of technology to relevant market sizes (supplies of factors of production working with these technologies and consumer demand), the price of other inputs into the production process (e.g., natural resources used in different sectors), markups, and regulations. The implications of the framework are broadly in line with a growing body of empirical work, especially from sectors such as energy, health care, agriculture, modern automation, and traditional industrial technologies.

More importantly for my focus here, the framework highlights several factors that can lead to *systematic* misalignment between market incentives and social objectives. First, some technologies generate negative externalities, while alternative paradigms aimed at performing similar production tasks may avoid these negative effects or even create positive externalities.⁴ The leading current example of this phenomenon is in the energy and transport sectors, where fossil-fuel-based energy creates carbon emissions and environmental damages, while renewables avoid such emissions. When the market does not price these damages, equilibrium innovation will be excessively directed toward fossil fuel technologies. I argue that similar issues arise in health care, where some technologies—for instance, those targeting prevention—may have greater social benefits than those aimed at high-tech procedures for late-stage cures. I will also suggest that the direction of technology may be distorted toward automation and away from worker-complementary technologies because labor market imperfections create a wedge between the social cost of labor and the equilibrium wage.

Second, different sectors often have different markups, and I show that equilibrium incentives will be excessively biased toward higher-markup sectors and technologies.⁵ Health care illustrates

³Existing evidence also suggests that although government encouragement to invest in high-tech sectors can have major benefits (e.g., Gruber and Johnson 2019; Moretti, Steinwender, and Van Reenen 2019), top-down research often generates extensive distortions. For example, Howell et al. (2021) show that traditional Defense Department research contracts have become less effective over time, while those that provide more open-ended support for new areas create more successful innovations. Similarly, convincing evidence of productivity benefits from industrial policy comes from settings in which such policy supported broad sectors, such as heavy and chemical industries in South Korea and Finland (Lane 2021; Mitrunen 2019). Additionally, Branstetter, Li, and Ren (2022) show that recent Chinese industrial policy has not been successful in increasing firm productivity, while Acemoglu, Yang, and Zhou (2023) provide evidence that top-down Chinese academic incentives have led to significant distortions in the direction of research.

⁴These positive externalities may also be on future research—for example, with some areas creating more substantial knowledge gains upon which future innovations can build. Another example of negative externalities would be “defensive innovations” undertaken by incumbents in order to prevent rivals from increasing their market share.

⁵High markups encourage more innovation effort but also simultaneously reduce the utilization of a technology. This latter effect implies that expanding the production level of high-markup technologies is also socially valuable. Yet

this phenomenon, for curative technologies appear to have higher markups than preventative ones.⁶

Third, a variety of social forces may favor one paradigm ahead of alternatives. For example, the research community may value certain types of breakthroughs more than others because they are viewed as the more exciting research area or because they are more useful for building a scientific reputation. One possible illustration may be from modern digital technologies, where many researchers believe that the most important and coveted advances are those that enable algorithms to reach human parity in a range of tasks. This perspective then creates greater incentives to work on automation rather than other paradigms aimed at more human-complementary tools.⁷

Fourth, when different technologies create distinct distributional effects and society cares about inequality (for either direct or indirect reasons related to political economy), the market will fail to internalize these additional considerations.

Fifth, the direction of innovation may be distorted because of coordination failures. For example, there may be insufficient diversity in research investments, or firms and innovators may fail to coordinate on more productive alternative paradigms, as I discuss briefly below.

While, in practice, all five of these effects are likely important, I focus on the first two for two related reasons. First, these two channels can be, in principle, quantified by measuring markups or social benefits/externalities and are thus better candidates for “systematic distortions” in the direction of innovation. In the last part of the paper, I make a preliminary attempt at this type of quantification. Second, because

these are quantifiable distortions, correcting them does not require government agencies to have superior information or an ability to “pick winners.” By comparison, it is more difficult to objectively determine whether an untried paradigm would be more successful or whether the research community’s enthusiasm for a specific topic is a “fad” leading to excessive concentration of innovative effort.

I reestimate the empirical models from three studies in order to obtain some of the key parameters of my framework in these three different settings. These studies are Acemoglu and Restrepo (2022) for research directed at automation technologies; Acemoglu et al. (2023) for medical research directed toward different types of diseases; and Aghion et al. (2016) for fossil-fuel-based and cleaner innovations in the automobile industry. In each case, I identify the elasticity of substitution between different technologies and the degree to which past advances in one field create a relative advantage for the same field in the future. These two parameters are critical both for the equilibrium response of the direction of innovation to factor supplies, prices, and policies and for the divergence of equilibrium allocations from socially optimal choices. I combine these numbers with estimates of the social costs/benefits of different technologies and markups to assess how distorted equilibrium technology choices are, and the welfare gains from redirecting innovation. In each case, I provide suggestive evidence that innovation distortions and their welfare effects are sizable.

Related Literature.—This paper builds on and extends the literature on directed technological change. The first explicit discussion of this topic is in Hicks’s (1932) argument that a high price for a factor induces technological change targeted at economizing on that factor. The induced technology literature of the 1960s explored whether technological change would be Harrod neutral (purely labor augmenting) as typically imposed in neoclassical growth models (e.g., Kennedy 1964; Ahmad 1966; Samuelson 1965; Drandakis and Phelps 1966). But this literature relied on ad hoc rules to determine the direction of technology, and the exact form of these rules had defining effects on their results.

An early empirical investigation of these issues was Habakkuk’s (1962) seminal study of

this force is dominated by the former effect, and thus the equilibrium involves excessive innovation effort devoted to high-markup technologies.

⁶One telling example is emphasized in Howard et al. (2015): the melanoma drug Yervoy, approved in 2011, was marketed at the price of \$120,000 for a four-dose treatment by the pharmaceutical company Bristol Myers Squibb. It extends life by about four months. I provide direct evidence on these markup differences in Section IV.

⁷See Acemoglu and Restrepo (2020b); Acemoglu, Jordan, and Weyl (2021); Brynjolfsson (2022); and Acemoglu and Johnson (2023) for the argument that the general incentives in the artificial intelligence (AI) community and the agenda spearheaded by Alan Turing are creating an excessive focus on using AI technologies for automation.

American technology in the nineteenth century. Habakkuk (1962) argued that the direction and speed of American technology was shaped by a desire to economize on scarce skilled labor in the country. Allen (2009) similarly suggested that Great Britain was the first country to industrialize because British labor was more expensive than labor in other European economies and in China.

The more recent literature on the direction of technological change has developed models with explicit R&D decisions targeted at different sectors and monopolistic profits from new technologies, shaping the composition of innovation. The first two papers within this area, Acemoglu (1998) and Kiley (1999), investigated the reasons why recent industrial technologies have often been skill biased and why this skill bias may have accelerated starting in the 1980s, concurrently with the large increase in the supply of educated workers in the United States and other industrialized nations.

Acemoglu and Zilibotti (2001) develop a model of directed technological change in a multicountry setup, whereby technology choices in advanced economies are directed to their own needs, creating a form of “inappropriate technology” from the viewpoint of less developed nations with different factor endowments. Broader cross-country implications of directed technological change are studied in Gancia and Zilibotti (2005). Acemoglu (2003b) and Thoenig and Verdier (2003) explore how trade openness impacts the endogenous skill bias of technology. Acemoglu (2003a) and Jones (2005) investigate why an endogenous direction of technology can lead to Harrod-neutral advances and thus balanced growth as in textbook neoclassical models. Acemoglu and Linn (2004) and Costinot et al. (2019) study how demographic changes that alter the future market sizes of different types of medical technologies impact the direction of innovation, while in Acemoglu (2010) I present a formalization of the Habakkuk-Allen hypothesis, where labor scarcity can be a spur to faster economic growth. Bovenberg and Smulders (1995); Goulder and Schneider (1999); Di Maria and Valente (2008); Grimaud and Rouge (2008); Acemoglu et al. (2012); Rodrik (2014); Acemoglu, Akcigit, and Kerr (2016); and Hémous (2016), among others, discuss the balance between clean and dirty technologies and possible corrective policies in the presence of

environmental externalities. Acemoglu and Restrepo (2018, 2022) and Hémous and Olsen (2022) explore the endogenous choice between automation and other types of technologies.

The model I present here generalizes Acemoglu (2002) in two important directions. First, I extend the baseline framework to perform comparative statics with respect to input prices, markups, and externalities. Second, to the best of my knowledge, I undertake the first general analysis of the efficiency of the direction of technology within this framework, though Acemoglu et al. (2012) provide a characterization of optimal policies to restore efficiency in a model of endogenous technology and carbon emissions.

Also closely related to this paper are a series of works that point out why the equilibrium direction of technology may be inefficient. Acemoglu (2012) shows that equilibrium innovation is often insufficiently diverse, investing too much in one of two alternative technologies. This innovation pattern then leaves the economy vulnerable to shifts in underlying conditions or blockages in existing technological paradigms. Acemoglu et al. (2016) propose a model in which the distribution of innovation between firms of different sizes and ages is distorted.

Akcigit, Hanley, and Serrano-Velarde (2021) distinguish between fundamental and applied research and argue that the former generates more knowledge spillovers. The paper provides empirical evidence and a quantitative evaluation of this source of inefficiency. Distortions in the direction of technology resulting from different knowledge spillovers and congestion effects are also explored in Dechezleprêtre, Martin, and Mohnen (2014); Hopenhayn and Squintani (2021); and Martin and Verhoeven (2022), and are present in models of international technology diffusion as well (e.g., Grossman and Helpman 1993; Coe and Helpman 1995). Acemoglu, Akcigit, and Kerr (2016) show that new ideas in some fields matter more for subsequent innovation than others. These differential knowledge spillovers are complementary to the distortions emphasized in this paper.

Another related literature focuses on the choice between different technological paradigms and the possibility of inefficient lock-in (e.g., Dosi 1982 and Arthur 1989, as well as recent work by Acemoglu and Lensman 2023).

The rest of the paper is organized as follows. Section I presents the basic framework of directed technological change, first in a static and then a dynamic setting, and characterizes equilibrium innovation. Section II compares the equilibrium allocation and the types of technologies developed to those that are socially optimal. Section III reviews several studies that provide evidence on the effects of market size, prices, markups, and policies on the direction of innovation. Section IV focuses on a few of these studies to obtain estimates of the key parameters of the framework. It then combines these parameters with numbers on markups and externalities to present a first evaluation of the extent of technology distortions and welfare gains from correcting them. Section V contains concluding comments, while the online Appendix presents additional derivations, results, and details left out of the main text.

I. A Simple Model of Directed Technology

In this section, I provide a simple, two-sector model of directed technology. For simplicity, I start with a static setting and then outline the dynamic version, which is similar to the setup in Acemoglu (2002).

A. Static Environment

The economy is static and inhabited by a representative household with preferences given by

$$(1) \quad U = \ln C + \ln E,$$

where C denotes consumption, while E is an externality term, specified below. The log functional form is adopted to maximize the similarity with the infinite-horizon version of this model. There are three different types of labor (two of them working in the two sectors, plus scientists). As is standard, I assume that labor income from all these types of labor accrues to the representative household.

The unique final good is produced with the production function

$$(2) \quad Y = \left[\gamma_1 Y_1^{\frac{\varepsilon-1}{\varepsilon}} + \gamma_2 Y_2^{\frac{\varepsilon-1}{\varepsilon}} \right]^{\frac{\varepsilon}{\varepsilon-1}},$$

where Y_1 and Y_2 denote the output levels of the two intermediate products, which are themselves

produced using two different types of technologies (clean versus dirty, preventative versus curative, worker friendly versus automation, and so forth). Their production functions are given by

$$(3) \quad Y_j = X_j^\alpha R_j^{1-\alpha}$$

for $j \in \{1, 2\}$, where X_j denotes a variable input aggregate and R_j is a resource input, with exogenous price q_j^R .

The variable input is in turn produced with the following production function:

$$(4) \quad X_j = \left(\int_0^{N_j} x_j(\nu)^{1-\beta} d\nu \right) \tilde{L}_j^\beta,$$

where \tilde{L}_j is a specialized factor employed only in sector j . For example, \tilde{L}_2 could correspond to skilled (or college-educated) labor, while \tilde{L}_1 could be unskilled labor.⁸ In addition, $x_j(\nu)$ denotes the quantity of the different machine varieties used in the production of intermediate good $j \in \{1, 2\}$ and $\beta \in (0, 1)$. In this formulation, $[0, N_j]$ denotes the range of machines used in the production of $j \in \{1, 2\}$ and captures how advanced the technology for intermediate good j is. Once invented, each machine can be produced at the fixed marginal cost $\psi > 0$ in terms of the final good.

In what follows, I assume that all labor is supplied inelastically:

$$\tilde{L}_j = L_j \text{ for } j \in \{1, 2\}.$$

I model the *innovation possibilities frontier*, which specifies how new machine varieties are invented, by assuming that new ideas (or new machine varieties) are created by scientists. Specifically, the technology for creating new machine varieties is assumed to take the following static form:

$$(5) \quad N_j = \tilde{\eta}_j \phi(S_j) S_j,$$

where $\tilde{\eta}_j > 0$, S_j is the number of scientists assigned to technology j , and $\phi(S_j) = S_j^{\delta/1-\delta}$ with $\delta \in [0, 1)$. When $\delta > 0$, the innovation possibilities frontier features increasing returns

⁸In some applications, these could be the same labor allocated to the two sectors. In that case, L_1 and L_2 would be endogenous and satisfy a single market-clearing constraint, $L_1 + L_2 \leq \bar{L}$. Equilibrium would then require their earnings in the two sectors to be equalized—that is, $w_2/w_1 = 1$.

to scale at the sectoral level. In the dynamic model, these increasing returns will take the form of *path dependence*, meaning that past advances in the technology of a sector will make further advances in the same sector easier. Scientists take the behavior of other scientists—and thus the value of the ϕ function—as given. The case of $\delta = 0$ corresponds to the useful benchmark in which there are no increasing returns to scale and scientists have a constant productivity in each sector.

Scientists that innovate and create varieties of machines become the owners of the technology monopolists that sell those varieties. This means that when a scientist invents a new machine for sector $j \in \{1, 2\}$, she will be able to make a profit π_j , which I characterize below.

The total number of scientists is fixed, so market clearing for scientists yields⁹

$$(6) \quad S_1 + S_2 = \bar{S}.$$

To determine the profit levels from new machines for the two sectors, I adopt a simple market structure where each sector is subject to a fringe of competitive firms that can imitate and produce every machine but do so less efficiently. This forces a limit price in each sector, given by

$$(7) \quad q_j = (1 + \mu_j)\psi,$$

where $\mu_j \in \left(0, \frac{\beta}{1-\beta}\right]$. This formulation provides a tractable form in which mark-ups are potentially different between the two sectors.

Finally, I assume that the externality term in (1) takes a simple form given by

$$(8) \quad E = e^{-\sum_{j \in \{1,2\}} \tilde{\tau}_j \ln N_j},$$

⁹Formally, in online Appendix A, I suppose that each scientist has mass $\mathfrak{s} > 0$ and then consider the limit where $\mathfrak{s} \rightarrow 0$. This only matters in ensuring that deviations are well defined in the presence of externalities. Additionally, note that the assumption that scientists take the value of $\phi(S_j)$ as given does not matter for the results because with a fixed supply of scientists and the iso-elastic form of the ϕ function, the allocation of scientists between the two sectors is the same even if scientists form consortia that internalize the positive externalities they create on other scientists working in the same field. Finally, it is straightforward to make the total supply of scientists endogenous to the income that they derive from innovation, but I will not do so in this paper.

where $\tilde{\tau}_j \geq 0$ represents a negative externality from technology $j \in \{1, 2\}$ (or, if there are positive externalities, then $\tilde{\tau}_j < 0$). The assumption that the negative externalities originate from the level of technology is adopted for simplicity. Because these externalities do not impact market prices and are ignored by scientists and firms, they will play no role in the equilibrium allocation but will have a major impact on the efficiency of the equilibrium.¹⁰

An equilibrium in this environment is defined as an allocation in which both the final good sector and the two intermediate sectors minimize costs, technology monopolists maximize profits by setting the limit price given in (7), scientists maximize their income by choosing which sector to innovate in, and all markets clear. I am particularly interested in the equilibrium level of relative technology, denoted by n^{EQ} (where $n \equiv N_2/N_1$).

B. Static Equilibrium

Cost-minimizing demands for machines and resources can be computed from the maximization problem

$$(9) \quad \max_{\{x_j(\nu), L_j, R_j\}} p_j \left(\int_0^{N_j} x_j(\nu)^{1-\beta} d\nu \cdot L_j^\beta \right)^\alpha R_j^{1-\alpha} - \int_0^{N_j} q_j(\nu) x_j(\nu) d\nu - w_j L_j - q_j^R R_j,$$

where w_j is the price (wage) of factor $j \in \{1, 2\}$. Combining (7) with the expressions for machine and resource demands (provided in online Appendix A), we obtain technology monopolists' profits as

$$(10) \quad \pi_j(\nu) = \mu_j \psi x_j(\nu) = \mu_j \psi \left[\left(p_j \left(\frac{(1-\beta)\alpha}{(1+\mu_j)\psi} \right)^\alpha \left(\frac{1-\alpha}{q_j^R} \right)^{1-\alpha} \right)^{\frac{1}{\alpha\beta}} L_j \right] \equiv \pi_j,$$

¹⁰If, instead, the externalities were from the production or consumption levels of the intermediates, technology would have additional indirect effects working through changes in equilibrium prices. The simplification enables me to remove these indirect effects.

where the square-bracketed expression in the second line is $x_j(\nu) \equiv x_j$ and the last equality defines the equilibrium flow profits for the two sectors, π_j , which, as claimed above, is identical for all machines used in sector $j = 1, 2$.

Setting the final product as the numeraire, the cost-minimization condition for the final good sector implies

$$(11) \quad p_j = \gamma_j \left(\frac{Y_j}{\bar{Y}} \right)^{-\frac{1}{\varepsilon}}.$$

Combining this expression with (3), (4), and the expressions for machine and resource demands in online Appendix A, we obtain

$$(12) \quad p \equiv \frac{p_2}{p_1} = \left(\frac{\gamma_2}{\gamma_1} \right)^{\frac{\alpha\beta\varepsilon}{\sigma}} \left(\frac{1 + \mu_2}{1 + \mu_1} \right)^{\frac{\alpha(1-\beta)}{\sigma}} \\ \times \left(\frac{q_2^R}{q_1^R} \right)^{\frac{1-\alpha}{\sigma}} \left(\frac{N_2}{N_1} \right)^{-\frac{\alpha\beta}{\sigma}} \left(\frac{L_2}{L_1} \right)^{-\frac{\alpha\beta}{\sigma}},$$

where $\sigma \equiv \alpha\beta\varepsilon + 1 - \alpha\beta$ is the *derived* elasticity of substitution between the two types of labor. Intuitively, the relative price of a sector's product is decreasing in the technology and labor supply to the sector, since these tend to expand its output. In addition, higher resource prices and markups increase a sector's relative price.

The price levels are then obtained by combining this relative price equation with the ideal price condition, which uses the fact that the final good is the numeraire:

$$(13) \quad \left[\gamma_2^\varepsilon p_2^{1-\varepsilon} + \gamma_1^\varepsilon p_1^{1-\varepsilon} \right]^{\frac{1}{1-\varepsilon}} = 1.$$

Factor prices are equal to the value of the marginal product of the relevant factor: $w_j = \alpha\beta p_j Y_j / L_j$ for $j = 1, 2$. Using this equation and (12), the relative wage of the two types of labor can be derived as

$$(14) \quad \frac{w_2}{w_1} = \left(\frac{\gamma_2}{\gamma_1} \right)^{\frac{\varepsilon}{\sigma}} \left(\frac{1 + \mu_2}{1 + \mu_1} \right)^{-\frac{(1-\beta)(\sigma-1)}{\beta\sigma}} \\ \times \left(\frac{q_2^R}{q_1^R} \right)^{-\frac{(1-\alpha)(\sigma-1)}{\alpha\beta\sigma}} \left(\frac{N_2}{N_1} \right)^{\frac{\sigma-1}{\sigma}} \left(\frac{L_2}{L_1} \right)^{-\frac{1}{\sigma}}.$$

This expression confirms that σ is indeed the elasticity of substitution between the two types of labor. Equation (14) additionally shows that N_2/N_1 plays the role of relative factor-augmenting technological change. We can see that $(\sigma - 1)/\sigma$ also regulates the impact of relative technology N_2/N_1 , markups, and resource prices, since the net effect of these economic quantities depends on whether they affect the production level of an intermediate good by more or less than its price.

In an interior equilibrium in which research is directed to both technologies, scientists should make the same profits from improving the technology for either sector. Recalling that the productivity of a scientist when she works on technology j is $\tilde{\eta}_j \phi(S_j) = \tilde{\eta}_j S_j^{\delta/1-\delta}$, an interior equilibrium must satisfy $\tilde{\eta}_1 S_1^{\delta/1-\delta} \pi_1 = \tilde{\eta}_2 S_2^{\delta/1-\delta} \pi_2$. Inverting $\phi(S_j)$, we have $N_j = \tilde{\eta}_j S_j^{1/1-\delta}$, or $S_j = (N_j/\tilde{\eta}_j)^{1-\delta}$. Substituting this into the condition for an interior equilibrium and defining $\eta_j \equiv \tilde{\eta}_j^{1-\delta}$, we have

$$(15) \quad \eta_1 N_1^\delta \pi_1 = \eta_2 N_2^\delta \pi_2.$$

Combining this equation with (10) and (12), we obtain

$$(16) \quad n^{EQ} = \left[\frac{\eta_2}{\eta_1} \left(\frac{\gamma_2}{\gamma_1} \right)^{\frac{\varepsilon}{\sigma}} \frac{\mu_2}{\mu_1} \left(\frac{1 + \mu_2}{1 + \mu_1} \right)^{-\frac{\sigma - (1-\beta)}{\beta\sigma}} \right. \\ \left. \times \left(\frac{q_2^R}{q_1^R} \right)^{-\frac{(\sigma-1)(1-\alpha)}{\alpha\beta\sigma}} \left(\frac{L_2}{L_1} \right)^{\frac{\sigma-1}{\sigma}} \right]^{\frac{\sigma}{1-\delta\sigma}}.$$

Equation (16) links the equilibrium technology ratio between the two sectors to parameters of the final good production function, the innovation possibilities frontier, resource prices, markups, and the relative supplies of factors employed in the two sectors. For example, focusing on the case where $\delta\sigma < 1$, n^{EQ} is increasing in L_2/L_1 if and only if $\sigma > 1$, as I discuss in greater detail below.

The first proposition follows from this discussion and equation (16), and the uniqueness of equilibrium is established in online Appendix A.

PROPOSITION 1: *Suppose that $\delta < 1/\sigma$. Then there exists a unique equilibrium in which the relative technology ratio is given by (16).*

The comparative statics of the direction of technology in this unique equilibrium are provided readily by equation (16) and will be discussed after I present the dynamic version of this environment in the next subsection.

To understand the role of the condition $\delta < 1/\sigma$, note that the stabilizing economic force in this model is the lower price of the intermediate good that is technologically more advanced, as shown by (12), and this force is stronger when σ is lower. The destabilizing force, on the other hand, is the extent of increasing returns, δ . When $\delta < 1/\sigma$, the sector that is further ahead technologically faces sufficiently lower returns from innovation, and this ensures the existence and uniqueness of the interior equilibrium.

In contrast, when $\delta > 1/\sigma$, the degree of increasing returns to scale in research is sufficiently strong that there does not exist an equilibrium in which research is directed toward both sectors. I show in online Appendix A that in this case there are two corner equilibria—all scientists working in sector 1 or all scientists working in sector 2. This discussion also illustrates why the comparative statics of (16) are only relevant when $\delta < 1/\sigma$.

C. Dynamic Environment

I now present the dynamic version of this economy, for brevity emphasizing only the elements that are different from the static setup. Suppose that time is continuous and runs to infinity. There is an infinitely lived representative household with preferences given by $U(0) = \int_0^\infty e^{-\rho t} U(t) dt$, with $U(t) = \ln C(t) + \ln E(t)$, where $C(t)$ denotes consumption at time t , $E(t)$ is the externality term as in the text, and ρ is the discount rate of the representative household. Analogously with (8), we have $E(t) = e^{-\sum_{j \in \{1,2\}} \bar{\eta}_j \ln N_j(t)}$.

All of the equilibrium conditions derived in the static model now apply, except that they should be indexed by time. The main difference is the innovation possibilities frontier, which takes the form

$$(17) \quad \dot{N}_j(t) = \eta_j N_j(t)^{(1+\delta)/2} \times N_{\sim j}(t)^{(1-\delta)/2} S_j(t),$$

where $S_j(t)$ denotes scientists working for innovation in technology j at time t and the

$N_j(t)^{(1+\delta)/2}$ term captures path dependence in innovation from one's own sector, while $N_j(t)^{(1-\delta)/2}$ is the contribution of the technology of the other sector. This innovation possibilities frontier is the dynamic analogue of (5). Notice the difference from the increasing returns to scale in (5): in the dynamic case, there is a form of increasing returns to scale, but it is realized over time. This is the reason I refer to δ as the degree of *path dependence*—when $\delta > 0$, once a sector is technologically ahead of the other one, it becomes more productive in generating new innovations.

The total number of scientists is again fixed, so $S_1(t) + S_2(t) = \bar{S}$ at all t . Scientists who innovate and create new varieties now become the perpetual owners of the technology monopolists that sell those varieties. Suppose also that resource prices, the q_j^R 's, are constant, which implies that profits from technology $j = 1, 2$ in this dynamic environment are constant and are still given by π_j in equation (10).

I first focus on an interior balanced growth path (BGP) in which $n(t) \equiv N_2(t)/N_1(t)$ is constant and thus scientists work on both technologies. This requires

$$\eta_1 \pi_1 N_1(t)^\delta = \eta_2 \pi_2 N_2(t)^\delta \text{ for all } t.$$

This condition is identical to (15) in the previous subsection. Hence, the BGP technology ratio in this dynamic model is identical to the equilibrium technology ratio in the static model.

PROPOSITION 2: *There exists a unique BGP, where the equilibrium direction of technology is given by equation (16).*

The fact that the unique BGP ratio coincides with the static equilibrium technology ratio is because of the way in which the static model was set up to mimic the insights of the dynamic framework.¹¹ While the BGP here coincides with the equilibrium of the static model, the full equilibrium path of the dynamic model leads to somewhat different results, as explained in the next proposition.

¹¹ The existence of a unique BGP is a consequence of the simplifying functional form assumptions. In general, as discussed in Acemoglu (2007), multiple equilibria are possible. But given my focus here, uniqueness enables me to focus on issues of distorted technology more directly.

PROPOSITION 3: *If $\delta < 1/\sigma$, the unique interior BGP (given by (16)) is globally (saddle-path) stable. In particular, starting from any initial conditions, the economy tends to this interior BGP. Moreover, the unique dynamic equilibrium allocates all scientists to the sector that is relatively behind (compared to the BGP). As a result, the BGP is reached in finite time.*

If $\delta > 1/\sigma$, then the interior BGP is unstable, and starting from almost all initial conditions the economy limits to an allocation in which only one of the two technologies advances.

This proposition clarifies why the case with $\delta < 1/\sigma$ is the focal one in the dynamic economy as well, and the intuition for this condition is similar: the stabilizing force via relative price changes should be stronger than the destabilizing force due to path dependence.

When $\delta > 1/\sigma$, the equilibrium in the dynamic model is still unique (in contrast to the static model, where there were multiple equilibria), but now the relative technology level identified by condition (16) corresponds to an unstable BGP and the economy will never converge to it. Rather, the equilibrium allocation will limit to one of the two corner BGPs, where the economy has a constant growth rate, driven by research in only one of the two technologies.

D. Some Properties of Equilibrium Technology Choices

I now review some properties of equilibrium direction of technology (using either the static equilibrium or the BGP of the dynamic equilibrium). This discussion will be brief because most of this material is familiar from previous work and is not my main focus here, though recognizing these comparative statics helps build intuition about the workings of the model.

Relative Supply Effects.—The direction of technology is determined by the relative supply of labor used with the two types of technologies, L_1 and L_2 . As in Acemoglu (1998, 2002), the implications of relative supplies on the direction of technology depend on market size and price effects. Holding prices p_1 and p_2 fixed, greater relative supply of one type of labor expands the market size of the technology complementing that type of labor, and it further encourages the development of this

complementary technology (i.e., a higher L_2/L_1 increases n^{EQ}). However, in equilibrium, prices also adjust, and this creates a countervailing force. Whether this countervailing force is more powerful than the direct market size effect depends on the elasticity of substitution between the two types of labor. Specifically, when $\sigma > 1$, the market size effect dominates the price effect, and n^{EQ} is increasing in L_2/L_1 . In contrast, when $\sigma < 1$, the price effect is more powerful, and n^{EQ} is decreasing in L_2/L_1 .

Weak Bias of Technology.—As we have just seen, the impact of L_2/L_1 on n^{EQ} is ambiguous and depends on the elasticity of substitution between the two factors. Nevertheless, as emphasized in Acemoglu (2002), there is a general, unambiguous result about the bias of technology. A change in technology is said to be *biased toward a factor* if, holding all other variables constant, it increases the relative price of that factor. The main result that holds in this class of models is that an increase in L_2/L_1 always induces a change in technology that is (weakly) biased toward L_2 . For example, if the relative supply of college-educated workers increases, then technology becomes more skill biased. Intuitively, this is because, as equation (14) demonstrates, when $\sigma > 1$, a greater L_2/L_1 raises n^{EQ} , and in this case it is also a higher level of n^{EQ} that is biased toward type 2 workers. Conversely, when $\sigma < 1$, it is a *decrease* in n^{EQ} that is biased toward type 2 workers, and in this instance, higher L_2/L_1 leads to *lower* n^{EQ} . Hence, regardless of the exact value of the elasticity of substitution between the two factors, technology always (weakly) moves in a direction that is favorable to the more abundant factor. Among other things, this force might explain why aggregate technology has become more skill biased over the last eight decades, while the supply of skilled workers in the industrialized world has risen rapidly (Acemoglu 1998). Acemoglu (2007) shows that this weak bias result is more general and holds without any of the functional form assumptions imposed here, provided that some mild regularity conditions are satisfied.

Strong Bias of Technology.—By substituting the expression for n^{EQ} from (16) into (14), we obtain the long-run (endogenous technology)

relationship between relative supplies and relative wages as

$$(18) \quad \left(\frac{w_2}{w_1}\right)^{BGP} = \Gamma \left(\frac{L_2}{L_1}\right)^{\frac{\sigma-2+\delta}{1-\delta\sigma}},$$

with

$$\Gamma \equiv \left[\frac{\eta_2}{\eta_1} \left(\frac{\gamma_2}{\gamma_1}\right)^{\frac{\varepsilon(1-\delta)}{\sigma-1}} \left(\frac{\mu_2}{\mu_1}\right) \times \left(\frac{1+\mu_2}{1+\mu_1}\right)^{-\frac{1-\delta(1-\beta)}{\beta}} \left(\frac{q_2^R}{q_1^R}\right)^{-\frac{(1-\alpha)(1-\delta)}{\alpha\beta}} \right]^{\frac{\sigma-1}{1-\delta\sigma}},$$

where recall that $\delta\sigma < 1$. This equation implies that the relationship between relative wages and relative supplies is upward sloping when $\sigma > 2 - \delta$, exactly as in Acemoglu (1998, 2002). Intuitively, the condition $\sigma > 2 - \delta$ ensures that technology moves sufficiently in the direction of the factor that becomes more abundant. With this powerful change in the direction of technology, the demand for the more abundant factor increases so much that the overall consequence is to raise this factor's marginal product more than that of the less abundant factor. Consequently, the locus of long-run equilibria becomes upward sloping—greater relative supply translates into greater relative wage. Notice that when technology is fixed, relative demand curves are always downward sloping in this model (as in all models with price-taking firms). The upward-sloping demand curve is a consequence of technology's response to changes in relative supplies of factors. Acemoglu (2007) provides a version of the same result for more general technologies and also an analogue of this result for the wage level of a factor rather than its relative wage.

Resource Prices.—Equation (16) also clarifies that resource prices will have a major impact on the direction of technology. In particular, when $\sigma > 1$, an increase in q_2^R (relative to q_1^R) reduces n^{EQ} because higher resource prices for sector 2 make production, and thus the technologies being used in this sector, less profitable.

Effects of Markups.—Equation (16) further highlights a first-order effect of markups on the direction of technology, and under relatively

weak conditions, a higher μ_2 (holding μ_1 constant) increases n^{EQ} .¹²

Many of these theoretical implications receive empirical support, as I discuss in Section III.

II. Distorted Technology

This section compares the equilibrium and socially optimal technology choices and identifies several reasons why equilibrium technology choices will be distorted.

A. Socially Optimal Direction of Technology

I now consider the social planner's solution in the static environment (the same exercise for the dynamic setup is presented in online Appendix A). Differently from equilibrium incentives, the social planner takes into account the externalities that the two intermediates generate. Naturally, the planner also cares about the full income stream accruing to the representative household rather than just the monopoly profits.

In what follows, I further focus on the case in which the social planner cannot directly control prices and allocations—and thus will not be able to correct for externalities and markups by introducing Pigovian taxes/subsidies. This choice has three motivations. First, practical (information-related) or political constraints often prevent governments from removing monopoly markups or may even make it difficult to implement corrective taxes for externalities. Second, as discussed in Acemoglu et al. (2012), Pigovian taxes are not always sufficient by themselves to restore optimality when the direction of technology is endogenous.¹³

¹²The countervailing force here is that higher markups reduce output and, via this channel, increase prices. It is straightforward to verify that more research is directed to sector $j = 1, 2$ when its markup μ_j increases, provided that $\sigma\beta + \mu_j(1 - \sigma)(1 - \beta) > 0$. This condition is satisfied whenever $\sigma \leq 1$ or whenever μ_j is not too large.

¹³This is because in models with endogenous innovation, there are distortions both in the production sector (captured by the externalities targeted by Pigovian taxes) and in the allocation of research effort between different sectors (due to monopoly profits and knowledge externalities, such as the path dependence introduced above). As a result, optimal allocations should correct for both sets of distortions. For example, in the context of the energy sector, relying just on carbon taxes without actively redirecting technological

Third, this choice also enables me to clearly focus on the distortions created by the allocation of research effort and the welfare gains from eliminating these technology distortions (rather than the full welfare consequences of various microeconomic distortions).

Given these assumptions, the only choice of the social planner is the allocation of scientists between the two technologies. In practice, this can be achieved by targeted research subsidies or regulations, and here I assume that the planner directly controls this allocation. Hence, in the static environment, the planner's problem can be written as maximizing (1) by choosing S_1 and S_2 subject to (6) and the innovation possibilities frontier (5) and taking all other equilibrium relationships, in particular the price function (12), as given. This yields a simple maximization problem for the social planner:

$$\max_{S_1, S_2 \geq 0: S_1 + S_2 \leq \bar{S}} \ln Y[N_1, N_2] + \ln E[N_1, N_2]$$

subject to (5), (6), and (12). Taking the first-order conditions for this expression, noting that $d \ln N_j = dN_j/N_j$, and substituting for S_j in terms of N_j as in the equilibrium analysis, this necessary condition for an interior social optimum can be written as

$$(19) \quad \eta_1 \left[\frac{d \ln Y}{d \ln N_1} + \frac{d \ln E}{d \ln N_1} \right] = n^{-(1-\delta)} \eta_2 \left[\frac{d \ln Y}{d \ln N_2} + \frac{d \ln E}{d \ln N_2} \right].$$

Clearly, $d \ln E / d \ln N_j = -\tilde{\tau}_j$, and in online Appendix A I prove that $d \ln Y / d \ln N_j = \gamma_j^\varepsilon p_j^{1-\varepsilon}$. Moreover, defining $\tau_j \equiv \tilde{\tau}_j / (\gamma_j^\varepsilon p_j^{1-\varepsilon})$ as a price-adjusted externality, the first-order condition can be simplified to

$$\eta_1 \gamma_1^\varepsilon p_1^{1-\varepsilon} (1 - \tau_1) = \eta_2 p_2^{1-\varepsilon} \gamma_2^\varepsilon (1 - \tau_2) n^{-(1-\delta)}.$$

change away from fossil fuels would slow down the transition to clean energy and amplify its short-run costs.

We can then substitute from (12) and solve for the socially optimal ratio of technology between the two sectors, n^{SP} , as

$$(20) \quad n^{SP} = \left[\frac{\eta_2}{\eta_1} \left(\frac{\gamma_2}{\gamma_1} \right)^{\frac{\varepsilon}{\sigma}} \left(\frac{1 + \mu_2}{1 + \mu_1} \right)^{\frac{1-\beta}{\beta} \frac{1-\sigma}{\sigma}} \times \left(\frac{1 - \tau_2}{1 - \tau_1} \right) \left(\frac{q_2^R}{q_1^R} \right)^{-\frac{(\sigma-1)(1-\alpha)}{\alpha\beta\sigma}} \left(\frac{L_2}{L_1} \right)^{\frac{\sigma-1}{\sigma}} \right]^{\frac{\sigma}{1-\delta\sigma}}.$$

It is also useful to write the ratio of socially optimal and equilibrium technologies as

$$(21) \quad \frac{n^{SP}}{n^{EQ}} = \left[\left(\frac{\mu_2}{\mu_1} \right)^{-1} \left(\frac{1 + \mu_2}{1 + \mu_1} \right) \left(\frac{1 - \tau_2}{1 - \tau_1} \right) \right]^{\frac{\sigma}{1-\delta\sigma}}.$$

It can be verified that given τ_1 and τ_2 , a higher μ_2 always implies a lower n^{SP}/n^{EQ} .¹⁴ There are indirect effects from markups, but the overall impact from a higher sectoral markup is to distort technology toward that sector. Additionally, a higher τ_2 always implies a lower n^{SP} and n^{SP}/n^{EQ} because of the negative externalities. Finally, the impact of all of these factors on the extent of technology distortion is amplified by $\sigma/(1-\delta\sigma)$. This is because a higher elasticity of substitution between factors and a greater degree of increasing returns to scale (or path dependence) in innovation makes the equilibrium direction of technology more responsive to markups and the social planner's preferred direction more sensitive to externalities. The next proposition summarizes these results.

PROPOSITION 4: *Suppose that $\delta < 1/\sigma$. Then the social planner's problem has a unique solution given by (20). Greater externalities and higher markups in sector j imply that*

¹⁴ Recall that τ_1 and τ_2 are functions of prices, and in (20), they are evaluated at the relative technology level n^{SP} . However, the interpretation of (21) requires some caution. In particular, in writing this expression, we have to hold τ_1 and τ_2 fixed. Or, alternatively, when distortions are small, intermediate prices p_1 and p_2 under n^{EQ} and n^{SP} will be approximately the same, and thus a given level of $\tilde{\tau}_j$ will map to approximately the same level of τ_j .

equilibrium technology is excessively distorted toward sector j .

This proposition implies that the only sources of divergence between the equilibrium and the social planner's solution in this setting are due to markups and externalities. The social planner would like to move equilibrium technology away from sectors that have high markups and high negative externalities.

Equation (21) additionally implies that technology distortions can be quantified using four sets of quantities: markup differences, μ_2/μ_1 ; externality differences, $(1 - \tau_2)/(1 - \tau_1)$; the degree of increasing returns to scale, δ ; and the elasticity of substitution between the two types of labor used in the two sectors, σ (which is in turn a function of ε , α , and β).

As in the equilibrium characterization in Proposition 1, Proposition 4 focuses on the case where $\delta < 1/\sigma$. When this condition is violated, the social planner prefers all scientists to work only in one of the two sectors (see online Appendix A), whereas, as we have seen, all scientists working in either sector is an equilibrium.

Given the equilibrium characterization, one can also compute the welfare loss in equilibrium relative to the social optimum. Online Appendix A provides a first-order approximation to the change in welfare between the equilibrium and the social optimum, and I present estimates of this welfare loss in the context of the applications in Section IV.

B. Other Considerations

Before moving to an assessment of the quantitative extent of distorted technology in various applications, I comment on a few additional issues.

First, I have simplified the discussion by ignoring other sources of distortions in the direction of technology. One potentially important type of distortion originates from visions, beliefs, fads, and ideologies. For example, the private sector may come to believe that only one path of development of a scientific platform is feasible or may be gripped by a “technology fad.” These issues are discussed in Acemoglu and Restrepo (2020b) and Acemoglu and Johnson (2023) in the context of AI—arguing that the influence of dominant companies and certain research approaches developed in the

1950s and 1960s pushed the field too much toward automation-related applications. These considerations can be introduced in the current model in a reduced-form manner by assuming that the market's assessment of η_1 and η_2 are systematically biased away from the true values of these parameters. Alternatively, one of the sectors may offer greater reputation-building opportunities to researchers. The more interesting question, which is beyond the scope of the current paper, is how such misperceptions or distorted incentives arise and whether there could be systematic ways in which government regulation could detect and prevent them.

Second, for tractability's sake, I have assumed that the degree of increasing returns to scale, captured by the parameter δ , is the same in the two sectors. In practice, certain types of research—for example, those targeting a scientific breakthrough or the “research” rather than the “development” part of R&D—may generate more knowledge spillovers (e.g., Akcigit, Hanley, and Serrano-Velarde 2021). Such spillovers can also be introduced in our context, though measuring the exact extent of such externalities is challenging.

Third, policymakers may also wish to take into account distributional and other social effects.¹⁵ If society engages in costly fiscal redistribution in order to increase the incomes of certain groups (e.g., the unemployed, low-skill workers, and so forth), then we can think of technologies that directly increase these groups' productivity as generating first-order pecuniary externalities, which can again be captured by our τ parameters.

Fourth, there may be reasons why the market underinvests in diverse technologies, as argued in Acemoglu (2012). Specifically, when there are shifts in which technologies are appropriate in different time periods, the market economy may underinvest in having a diverse portfolio of technologies that can act as a stepping stone when the underlying environment changes.

¹⁵ Inequality generated by some technologies may create additional social problems (as argued, e.g., by Wilson 1996 and documented by Autor, Dorn, and Hanson 2019) or may erode support for democracy (as shown in Acemoglu et al. 2021). These considerations would constitute additional reasons for altering the direction of technology. Since these effects are harder to quantify, they fall beyond the scope of the current paper.

Finally, in richer models, there can be coordination failures whereby the market coordinates on or stays too long with an inferior technology (see Acemoglu and Lensman 2023 for recent work on this topic). Once again, quantifying distributional, diversity, and coordination effects is more challenging, and I leave these issues for future work as well.

III. Existing Empirical Evidence

In this section, I review several empirical papers from the areas of energy, health technologies, agriculture, modern automation technologies, and the introduction of new industrial machinery during the industrial revolution to provide an overview of a body of growing empirical evidence on how market sizes, resource prices, and policy impact the direction of technology. The available evidence generally supports the predictions of the theoretical framework of this paper.

A. Energy

There is a large and growing literature that shows the responsiveness of both energy-generation and energy-use technologies to resource prices. Newell, Jaffe, and Stavins (1999) studied the impact of energy prices on energy-saving innovations. The authors collected data on the cooling/heating capacity; energy flow; energy efficiency; and price of room air conditioners, central air conditioning units, and gas water heaters from the Sears, Roebuck and Company catalogs between 1958 and 1993. Their results show that higher energy prices have a significant impact on energy efficiency; the models offered to consumers became more energy efficient when resource costs rose. The authors also present some evidence that energy standard regulations had a similar effect for room air conditioners. Consistent with the idea that there is a strong trade-off between different types of technologies, the authors additionally show that energy efficiency adjustments are associated with higher prices, and, in fact, they do not find significant effects on the overall amount of technological change. Hence, this study suggests that the direction of technology may be more responsive to resource prices than the overall amount of technological change, which is consistent with the framework presented here when $\sigma > 1$.

Popp (2002) studies US patent and citations data from 1970 to 1994. He establishes a robust association between energy prices and energy-efficient innovations. He also shows a significant role of the knowledge base, reminiscent of path dependence in the innovation possibilities frontier above.

Aghion et al. (2016) provide additional evidence consistent with these patterns. These authors build a firm-level dataset of automobile-related patents across 80 countries and classify these innovations into dirty and clean technologies—for example, internal combustion engine versus hybrid and electric vehicles. They show that higher fuel prices induced by carbon taxes lead to more clean and less dirty innovations in the automobile industry. They also estimate statistically significant path dependence. In Section IV, I use this study's data to present some related results as a basis of my quantitative exercise.

More recent work by Acemoglu et al. (2019) documents a relationship between natural gas prices, driven by the US shale gas boom, and overall green patenting (relative to either all patents, energy patents, or dirty patents). In particular, green patents surged when natural gas prices were high and then declined as the shale gas boom kicked in.

Overall, the evidence from the energy sector is fairly clear that resource prices have the expected impact on the direction of technology and that the direction of technology is possibly more responsive than the overall amount of innovation. There is also evidence of path dependence, whereby energy-efficient (or green) innovations build on a specific knowledge base that past innovations of this type have created.

B. Health and Medical Technologies

The direction of health-care and medical technologies appears to be highly responsive to market sizes, prices, and regulations, along the lines of the predictions of the framework presented here. Finkelstein's (2004) pioneering study focuses on several policy changes, expanding the market size for certain vaccines. Specifically, in 1991 the Center for Disease Control recommended that all infants be vaccinated against hepatitis B, while in 1993 Medicare began covering the full cost of influenza vaccination for Medicare recipients (without any copayments).

Finkelstein (2004) also looks at a 1986 reform indemnifying manufacturers from lawsuits from potential adverse reactions to childhood vaccines against polio; diphtheria-tetanus; measles, mumps, and rubella; and pertussis. Finkelstein (2004) estimates a 2.5-fold increase in the likelihood of clinical trials for the relevant vaccines following the policy-induced expansion of market size.¹⁶

Acemoglu and Linn (2004) focus more directly on the market size for new pharmaceuticals. They exploit variation originating from demographic change—for example, the baby boomer generation first creating demand for pharmaceuticals targeted at younger and middle-aged patients and later, as this cohort aged, for drugs targeting diseases for older patients. They find a powerful impact of market size on the introduction of new molecular entities, as well as the entry of new generics. Their baseline estimate suggests that a 1 percent increase in market size is associated with a 4 percent increase in new nongeneric drugs. In subsequent work, Acemoglu et al. (2006) provide suggestive evidence that Medicare induced an increase in pharmaceutical innovations targeted at the elderly. Costinot et al. (2019) provide similar evidence from a cross-country setting. These authors combine predictions about the direction of innovation with the home market effect (whereby countries specialize in and export products targeted at their home market) and document that countries invest more in and export drugs that have a greater demand among their home population.

More recent research by Acemoglu et al. (2023) assembles a comprehensive dataset of cross-country medical research and disease burdens impacting different countries. They estimate a strong association between the burden from a disease and research directed toward that disease. Below, I also present regression

evidence from the dataset compiled by these authors.

Finally, the only paper I am aware of that provides evidence relevant for the effects of markup differences is Budish, Roin, and Williams (2015). These authors observe that the US patent system, where protection is granted for a fixed term length, creates greater pecuniary incentives for late-stage cancer treatments relative to early-stage treatments and cancer prevention. They show that there is a powerful effect favoring late-stage treatments. This can be interpreted as a difference in markups between two (imperfectly substitutable) treatment modalities targeting the same underlying problem.

Overall, health-care and medical technologies provide ample evidence supporting the role of market size in the direction of innovation, and several of the studies show that policy-induced changes in market size have sizable effects on the direction of technology as well. There is additionally some evidence on the role of markups.

C. Agriculture

Early work by Hayami and Ruttan (1970) applied ideas from the induced innovation literature to agriculture, focusing on incentives for developing more or less capital-intensive agricultural methods. More recently, Moscona (2021) studied the long-run effects of the soil erosion and reduced soil productivity in the American Midwest following the Dust Bowl and found that agricultural innovation shifted toward more impacted crops in an apparent effort to make them more productive under the new soil conditions (see also Hornbeck 2012).

Related work by Moscona and Sastry (forthcoming) focuses on the changes in environmental conditions caused by global warming. Using granular data on new crops, these authors find that since the middle of the twentieth century, agricultural innovation has shifted toward crops that have greater exposure to extreme temperatures, and this has been driven by the types of technologies that are most related to environmental adaptation, such as new crop varieties that can be grown at higher temperatures by existing farmers. The innovation response in these two papers is consistent with the predictions of the framework here when price effects are more powerful than market size effects (that

¹⁶Finkelstein (2004) does not find an increase in medical trials and patents, which may be due to the fact that the relevant knowledge for additional rollout of these six vaccines already existed. We know from the more unique but sharper variation coming from the COVID-19 pandemic that entirely new vaccines, together with a new body of scientific knowledge, were created in response to the huge increase in the demand for vaccines against this novel virus (see, e.g., Zuckerman 2021).

is, if $\sigma < 1$). In contrast, if market size effects had been dominant ($\sigma > 1$), innovations should have been redirected toward crops that are less affected by the Dust Bowl and climate change, and less of the affected crops should have been cultivated. In contrast, it appears that because price effects are stronger, innovation attempted to make up for the reduced productivity of the affected crops.

D. Modern Automation Technologies

Following Acemoglu (2003a); Acemoglu and Restrepo (2018, 2022); and Hémous and Olsen (2022), the two factors here can be mapped to capital and labor to capture a reduced-form model of automation. A more microeconomic model of automation and task allocation between capital and labor, as in Zeira (1998), is developed in Acemoglu and Autor (2011) and Acemoglu and Restrepo (2018); see also Autor, Levy, and Murnane (2003).

Acemoglu and Restrepo (2022) provide a first empirical study of this issue, exploiting the facts that demographic change is taking place at different rates across countries and aging is expanding the demand for automation technologies by reducing the industrial workforce. This paper shows that demographic change has a large impact on the demand for robots and other automation technologies and then uses patents and exports of intermediate products to establish that countries with aging workforces file more patents for automation technologies and export more intermediates involved in automation. There is no similar impact for nonautomation technologies. This study thus establishes a powerful channel from the market size for automated production methods to both the innovation and adoption of automation technologies. I will use the data from this study in the next section as well.

More recent work by Dechezleprêtre et al. (2021) confirms and more deeply explores this relationship. The authors build a new firm-level dataset of automation innovations based on patent text and combine this with macroeconomic data across 41 countries. They estimate that higher wages for low-skill workers lead to more automation innovation. In addition, they exploit the Hartz labor market reforms in Germany, which led to lower protection for workers, and show that these reforms were associated with a reduction in automation innovations.

Finally, Clemens, Lewis, and Postel (2018) study the end of the Bracero Program, which brought about half a million Mexican immigrants to work in US farming. They find no discernible effect on agricultural wages and provide evidence that this is because the decline in the supply of unskilled labor induced the adoption of more mechanized production methods in US agriculture.

Overall, the evidence suggests that although many factors have impacted the development and introduction of modern automation techniques, a major boost has come from changes in the market size for these technologies, driven by declines in the supply of labor and corresponding higher wages, due to aging or changes in regulations.

E. British Industrial Revolution

There is also a small economic history literature providing evidence that at various turning points during the industrialization process, the direction of innovation was heavily shaped by market sizes and scarcity of labor and other inputs. In addition to Habakkuk's (1962) and Allen's (2009) work discussed in the Introduction, Hanlon (2015) studies the technological implications of the shortage of cotton in Britain created by the Union Navy's blockade of Southern shipping during the American Civil War. After the introduction of the cotton gin, the US South had become a major (slave-based) producer of cotton and the largest exporter of this crop to the expanding British industry. The blockade of Southern exports during the Civil War created an acute shortage of inputs to the British cotton industry, which in response turned to alternative cotton varieties grown in India (and, to a lesser extent, in Egypt and Brazil). The spinning technologies used at the time were adapted to the American cotton and could not be used on Indian and other varieties. Hanlon (2015) interprets this change as an expansion in the market size of these alternative cotton varieties, which should, according to the framework presented here, trigger a major expansion of complementary technologies. Hanlon (2015) documents that this is exactly what happened. There was a flurry of spinning innovations and patenting but no spike in other textile technologies, such as weaving, and no changes in nontextile patents (for which there was no major

change in market size). Moreover, by studying the variation in cotton prices, Hanlon (2015) shows that the induced-innovation response was large enough to cause the equivalent of the strong-bias result outlined above.

F. Inappropriate Technologies

Another implication of the framework presented here is that when a disproportionate share of innovative activity is concentrated in a few countries and researchers in these countries target their own economies' factor endowments and prices, then the global technology will be inappropriate to the needs of remaining countries, especially when their conditions are very different from those of innovative economies.

A recent important paper by Moscona and Sastry (2023) extends the framework in Acemoglu and Zilibotti (2001) and provides evidence that this inappropriate technology channel is present and quantitatively important. They establish that new crop varieties and seeds are developed to be resistant to pests and pathogens that are important in the United States and Western nations, while the major pests and pathogens in the rest of the world, though closely related, are distinct. As a result, the same agricultural technologies do not achieve high productivity in developing-world agriculture. Moscona and Sastry (2023) document that inappropriate agricultural technologies are generally not adopted in the developing world and, consequently, agricultural productivity remains low in these countries. They estimate that global agricultural output could be increased by about 58 percent if the direction of innovation were better targeted toward the agricultural conditions in less developed economies. Relatedly, Acemoglu et al. (2023) show that global medical research responds to disease burden in rich countries but not in poor countries, and Diao et al. (2021) provide evidence from Ethiopia and Tanzania that firms using Western, capital-intensive technologies are not increasing employment.

IV. Quantitative Evaluation

In this section, I discuss how the extent of technology distortions can be assessed in the leading applications considered here (automation, health, and energy). I first outline the econometric framework I use for estimating the

parameters σ and δ and discuss how markup and externality differences are calibrated. I then provide baseline estimates and a quantitative evaluation of distortions in the direction of innovation in these three sectors.

A. Econometric Framework

For automation technologies, I use the dataset on automation patents and demographic changes from Acemoglu and Restrepo (2022). For health care, I rely on the medical research and disease burdens dataset compiled by Acemoglu et al. (2023). For energy, I use the firm-level patenting and innovation dataset constructed by Aghion et al. (2016), who then combine this with information on policy-induced changes in the cost of gasoline.

In each case, I start from the dynamic innovation possibilities frontier for entity (country or firm) f and technology j :

$$(22) \quad \frac{\dot{N}_{fj}(t)}{N_{fj}(t)} = \xi_{fj}(t) \eta_{fj} \Gamma_j(t) \times N_{fj}(t)^{-\frac{1-\delta}{2}} N_{f \sim j}(t)^{\frac{1-\delta}{2}} S_{fj}(t),$$

which generalizes equation (17) by including a constant, $\eta_{fj} > 0$, parameterizing the productivity of entity f in technology area j ; a time effect, $\Gamma_j(t)$; and a random term, $\xi_{fj}(t)$, orthogonal to everything else. In addition, $S_{fj}(t)$ is a measure of research effort devoted by this entity to technology area j (e.g., the number of scientists allocated to this line), and $\delta \in [0, 1]$ again designates the degree of path dependence.

In this formulation, $\dot{N}_{fj}(t)$ is the flow of patents or innovations, while $N_{fj}(t)$ is the stock of patents/innovations, which is estimated following Cockburn and Griliches (1988) and Aghion et al. (2016) by assuming that the stock of knowledge embedded in past patents depreciates at some rate (and as in these papers, I set this rate of depreciation to 20 percent).

When there are only two types of technologies, as in the benchmark model, we can define $n_{ft} \equiv N_{f2}(t)/N_{f1}(t)$ as relative technology, take logs, and use the approximations $\Delta n_{ft} \approx \dot{N}_{f2}(t)/\dot{N}_{f1}(t)$ as we transition from continuous to discrete time to obtain

$$\ln\left(\frac{\Delta n_{ft}}{n_{ft}}\right) = \bar{\eta}_f + \bar{\gamma}_t - \rho \ln n_{ft} + s_{ft} + \bar{\xi}_{ft},$$

where I defined $\bar{\eta}_f \equiv \ln \eta_{ff} - \ln \eta_{f \sim j}$, $s_{ft} \equiv \ln S_{ff}(t) - \ln S_{f \sim j}(t)$, $\bar{\gamma}_t \equiv \ln \Gamma(t) - \ln \Gamma_{\sim j}(t)$, and $\bar{\xi}_{ft} \equiv \ln \xi_{ff}(t) - \ln \xi_{f \sim j}(t)$. I also set $\rho \equiv 1 - \delta$.

Suppose that we have a shifter/forcing variable at the country or firm level z_f (such as relative resource prices, market sizes, or policies) that alters the relative profitability of different technologies. Suppose also that the allocation of research effort between the two technologies at the firm level can be written as $s_{ft} = \chi \ln z_{ft} + \lambda \Delta \ln z_{ft}$.¹⁷ Substituting for this relationship, we arrive at the estimating equation

$$(23) \quad \ln \left(\frac{\Delta n_{ft}}{n_{ft}} \right) = \bar{\eta}_f + \bar{\gamma}_t - \rho \ln n_{ft} + \chi \ln z_{ft} + \lambda \Delta \ln z_{ft} + \bar{\xi}_{ft}.$$

The left-hand-side variable is the flow of relative patents or innovations normalized by stock of relative patents in the two technology areas. The forcing variable is also relative.

From estimates of (23), the key parameters necessary for quantifying the extent of distortions can be recovered. First, I set $\hat{\delta} = \max\{0, 1 - \hat{\rho}\}$, which imposes that the estimate for δ does not become negative in a few specifications in which $\hat{\rho}$ takes a value above 1. Moreover, long-run effects can be obtained from estimates of (23). In particular, in an interior BGP, we have $\dot{N}_{f2}(t)/N_{f2}(t) = \dot{N}_{f1}(t)/N_{f1}(t)$ in (22) and $s_{ft} = \chi \ln z_{ft}$, and thus the long-run relationship between relative technology and the forcing variable is $\ln n_{ft} = \text{constant} + \frac{\chi}{1 - \delta} \ln z_{ft}$. Estimated long-run effects can then be linked to the underlying parameters. Specifically, equation (16) implies that when the forcing variable is changes in market size, we have $\chi/(1 - \delta) = (\sigma - 1)/(1 - \delta\sigma)$, and for the case of energy, where the forcing variable is changes in energy prices, we have $\chi/(1 - \delta) = -(\sigma - 1)(1 - \alpha)/(\alpha\beta(1 - \delta\sigma))$.

The same economic relationships can be alternately estimated at the technology-field level by running the following regression

separately by field, which follows directly from (22):

$$(24) \quad \ln \left(\frac{\Delta N_{fjt}}{N_{fjt}} \right) = \eta_{jt} + \Gamma_{jt} - \frac{\rho}{2} \ln N_{fjt} + \chi \ln Z_{fjt} + \lambda \Delta \ln Z_{fjt} + \xi_{fjt}.$$

With only two research areas, this is equivalent to estimating (23). In the medical research regressions, there will be many more than just two areas, and hence focusing on this regression will be more meaningful.

B. Measuring Shares, Externalities, and Markups

Throughout, I use numbers from the US economy. The factor shares α and β are obtained from the Bureau of Economic Analysis Input-Output Use Tables. Table 1 provides a summary of these numbers for our three applications. For the automation application, I assume $\alpha = 1$ and take β to be the wage bill divided by the sum of the wage bill and expenditures on intermediate inputs for the manufacturing sector in 2012, which gives $\beta = 0.22$. For the health application, I again set $\alpha = 1$ and take β to be the wage bill divided by the sum of the wage bill and expenditures on intermediate inputs for the health-care sector in 2012, which gives $\beta = 0.55$. For the energy application, I take $1 - \alpha$ and $\alpha\beta$ to be, respectively, expenditures on material inputs divided by the sum of the wage bill and expenditures on intermediate and material inputs, and the wage bill divided by the sum of the wage bill and expenditures on intermediate and material inputs. This gives $\alpha = 0.86$ and $\beta = 0.32$.

The simplest method to measure the τ parameters is to start with existing estimates of externalities from certain economic activities. I then convert these externalities into the equivalent of $\bar{\tau}$ in our model, which is in consumption units (recall equation (1)). Throughout, I adopt the convention that the sector creating negative externalities is sector 2.

In the automation case, I follow Acemoglu, Manera, and Restrepo (2020), who interpret estimates of wage declines following job loss as proxying for quasi-rents that workers enjoy above and beyond the marginal cost of labor hours (and, thus, above the the social opportunity cost of employment). Hence, if automation

¹⁷This form follows, for example, when there are within-period diminishing returns or congestion effects in research (e.g., Acemoglu 1998).

technologies reduce employment, they create a negative pecuniary externality proportional to labor earnings. Assuming that for the target group (workers), consumption is approximately equal to labor earnings, this corresponds to $\tilde{\tau}$ in our model. I measure this externality by combining estimates from Acemoglu and Restrepo (2020a) on the effects of robots on employment with the average estimate of the extent of wage declines following job loss (15 percent, following the review of the literature in Acemoglu, Manera, and Restrepo 2020). The details are provided in online Appendix A. The resulting estimate is $\tilde{\tau}_2 = 0.07$, as shown in Table 1. As an alternative, more conservative estimate, I consider the case where only half of the workforce receives quasi-rents, which implies average quasi-rents of 7.5 percent and $\tilde{\tau}_2 = 0.03$.

In the health-care case, I interpret output Y as quality-adjusted life years (QALYs), which depend on expenditures and innovations in two broad categories: preventative versus curative technologies (used after the onset of disease). I allow these two types of technologies to have different markups and social benefits. This distinction and my approach are motivated by Kenkel (2000); Kremer and Snyder (2015); and Newhouse (2021). To measure social benefits, I use a sample of 71 new technologies that can be sorted into these two categories and then rely on existing estimates from the medical literature to obtain how much gain in QALYs is obtained per \$1 of total cost (up-front R&D spending plus per unit usage costs). These numbers indicate that there are fewer QALY gains from a dollar of spending in curative technologies than that in preventative technologies, and I interpret this shortfall as a negative externality from N_2 (curative) relative to N_1 (preventative). The baseline estimate of $\tilde{\tau}_2 = 0.37$ indicates that the QALY gains from the preventative category are about 60 percent larger than those from the curative category. The details of these technologies and the relevant calculations are provided in online Appendix C and in Appendix D (available upon request). In the baseline quantitative evaluation for health care, I set these externalities equal to zero and subsequently explore the implications of these additional distortions separately.¹⁸

¹⁸ Estimating the shortfall of QALYs from curative technologies should be viewed as an alternative to using markup

Broadly speaking, differences in externalities and markups between these two classes of technologies result from the fact that both the level of demand and the elasticity of demand for technologies that can be used after the onset of a disease are different from those for preventative ones because of individual incentives and insurance and public policy reimbursement rules (see Kremer and Snyder 2015 and Newhouse 2021).

In the energy application, I focus on negative externalities created by fossil fuel emissions. I use a worldwide social cost of carbon (CO_2) of $SCC = \$185$ per metric ton of carbon (in 2020 dollars), based on Rennert et al.'s (2022) estimate. For the baseline, I focus on US damages only since the other applications also ignore worldwide externalities. To convert this estimate to US-only damages, I use the ratio of US-to-worldwide damages from Resources for the Future's recent report (0.14), which gives $SCC \approx \$26$ per metric ton of carbon.¹⁹ These estimates are then converted into $\tilde{\tau}_2$ following the procedure described in online Appendix A. The resulting estimates are depicted in Table 1 as well.

Estimates for $\tilde{\tau}_2$ need to be converted to τ_2 . Recalling that $\tau_j = \tilde{\tau}_j / (\gamma_j^\varepsilon p_j^{1-\varepsilon})$ and also that $\gamma_j^\varepsilon p_j^{1-\varepsilon} = \gamma_j (Y_j/Y)^{\frac{\varepsilon-1}{\varepsilon}} = p_j Y_j/Y$, we have $\tau_2 = \left(\frac{p_2 Y_2}{p_1 Y_1 + p_2 Y_2} \right)^{-1} \tilde{\tau}_2$. Since estimates of $p_1 Y_1$ and $p_2 Y_2$ in the various approaches are likely to be imprecise (because of the difficulty of matching the conceptual categories here to data), I use the fact that this expression implies $\tau_2 \geq \tilde{\tau}_2$, and in the spirit of obtaining lower bounds on innovation distortions, I proxy τ_2 by $\tilde{\tau}_2$ in all three applications. As a result, my baseline estimates of technological externalities are $\tau_2 = 0.07$ in the automation case (or $\tau_2 = 0.03$ using the more conservative estimate of quasi-rents), $\tau_2 = 0.37$ in the health-care case, and $\tau_2 = 0.13$ in the energy case when I focus on social cost of carbon for the United

differences since differential markups will lead to different QALYs from preventative and curative technologies.

¹⁹ See <https://www.rff.org/publications/explainers/social-cost-carbon-101/>. Rennert et al.'s (2022) estimate is based on a discount rate of 2 percent. The Environmental Protection Agency's most recent preferred approach also suggests a similar social cost of carbon (\$190) based on a 2 percent discount rate. See <https://www.epa.gov/system/files/documents/2022-11/>.

TABLE 1—EXTERNALLY CALIBRATED PARAMETERS

Parameters	Description	Values
<i>Panel A. Automation</i>		
α	1 – Material Share	1
β	Labor Share divided by α	0.22
(μ_1, μ_2)	Markups (assumption)	(μ, μ)
(τ_1, τ_2)	Externality (quasi-rent = 15%)	(0,0.07)
(τ_1, τ_2)	Externality (quasi-rent = 7.5%)	(0,0.03)
<i>Panel B. Health</i>		
α	1 – Material Share	1
β	Labor Share divided by α	0.55
(μ_1, μ_2)	Markups (estimated)	(0.46, 1.70)
(τ_1, τ_2)	Externality (from QALYs)	(0,0.37)
<i>Panel C. Energy</i>		
α	1 – Material Share	0.86
β	Labor Share divided by α	0.32
(μ_1, μ_2)	Markups (assumption)	(μ, μ)
(τ_1, τ_2)	Externality (US damages)	(0,0.13)
(τ_1, τ_2)	Externality (world damages)	(0,0.94)

Notes: This table presents the values of the parameters used in the equilibrium and welfare analysis. Panel A is for the automation application, panel B for the health-care application, and panel C for the energy application. Material and labor shares are taken from the Bureau of Economic Analysis Use Table for 2012 (see text for details). Markups in panel B are computed from Compustat via the production function estimation method based on De Loecker, Eeckhout, and Unger (2020). Firm-level markups are aggregated to the technology level using firm cost shares. Online Appendix C provides more details and alternative estimates. Externalities are computed from wage declines following job loss, based on Acemoglu, Manera, and Restrepo (2020) in panel A; from the shortfall of QALY gains from curative technologies relative to preventative technologies (based on own calculations in online Appendix C) in panel B; and from Rennert et al.'s (2022) estimate of the social cost of CO₂, converted to US-equivalent damages and for worldwide damages in panel C. Further details are provided in the text, online Appendix A, and online Appendix C.

States and $\tau_2 = 0.94$ when taking full global damages into account.

Finally, I assume that markups are equal between the two technologies in the automation and energy applications. In the health-care application, I use data from health-related Compustat firms, sorted into the preventative versus curative technologies. I then use production function estimation or accounting data to obtain estimates

of markups for these two groups of firms.²⁰ The details and list of companies in each category are provided in online Appendix C. The baseline markup estimates, which follow De Loecker, Eeckhout, and Unger (2020), yield $\mu_1 = 0.46$ and $\mu_2 = 1.70$ for the period 1980–2016, as shown in Table 1.²¹ These markups are high, though broadly consistent with the numbers in De Loecker, Eeckhout, and Unger (2020). For example, their estimates of revenue-weighted markups for pharmaceutical and medicine manufacturing and for medical equipment and supplies manufacturing (the two four-digit industries most closely related to curative technologies) are, respectively, 3.41 and 2.14 (or cost-weighted markups of 2.97 and 1.91). These high numbers are also in line with the common view that certain medical procedures and pharmaceuticals are priced much above marginal cost in the United States, partly because of lack of regulation and partly because of employer-provided health insurance and Medicare reimbursement policies (see, e.g., Agnell 2004; Howard et al. 2015; Anderson, Hussey, and Petrosyan 2019; and Case and Deaton 2020).

²⁰Health-care firms in the preventative category include basic health providers, various companies specialized in diagnostics, and vaccine manufacturers, while curative ones include major pharmaceutical companies as well as high-tech medical equipment manufacturers. See Appendix D (available upon request) for a full list.

²¹Markup estimates from Compustat should be interpreted as simply suggestive since both capital and labor information from this dataset are subject to significant measurement error and it is impossible to separate the output and factor usage of multiproduct companies into different business lines. Moreover, as I show in online Appendix C, there are nontrivial fluctuations and trends in markup estimates. Nevertheless, online Appendix C also shows that using different methods for production function estimation yields very similar estimates. One conceptual issue, discussed in online Appendix C, is whether markups over marginal cost from variable inputs, as estimated by the production function approach, or accounting markups that subtract payments to quasi-fixed factors are more appropriate in this context. In particular, although accounting profits do not correspond to economic profits, they may be more informative about incentives for innovation and entry. Reassuringly, accounting markups for the two group of firms are comparable to our baseline estimates ($\mu_1 = 0.51$ and $\mu_2 = 1.35$), and using them instead yields broadly similar results, as also shown in online Appendix C. Finally, I experimented with applying the same methods to the energy sector as well, but because there are only a few firms that can be associated with clean technologies, these markups are unstable.

C. Estimates: Automation

In the context of automation, I focus on technologies targeting automation versus those that can broadly be thought to increase worker productivity. Columns 1 and 2 of Table 2 present estimates of equation (23) using five-year or ten-year patent counts sorted between automation and nonautomation technologies. Following Acemoglu and Restrepo (2022), I exploit medium-term, partially anticipated changes in demographics, which reduce the availability of labor to perform manual tasks across countries. I focus on anticipated (15- or 20-year) changes in the ratio of workers aged 56 and above to those between the ages of 25 and 55 as the measure of aging.²² The left-hand-side variable is the relative flow of automation patents compared to the relative stock of automation patents. On the right-hand side, I additionally control for GDP per capita, log population, and average years of schooling of the population at the beginning of the sample interacted with time dummies. These controls allow for flexible differential trends as a function of baseline characteristics. As in the original paper, these regressions are weighted by manufacturing employment in 1990, since patent data are significantly noisier for countries with smaller manufacturing employment levels. The sample period in this case is 1986–2015.

Throughout this table, I report heteroscedasticity-robust standard errors clustered to allow for serial correlation (at the country level in columns 1–4 and at the firm level in columns 5 and 6).

Column 1 in Table 2 depicts estimates from equation (23) for a full sample of 66 countries. The main parameters are estimated reasonably precisely. For example, the estimate of $\hat{\rho} = 0.77$ (standard error = 0.14) implies a value of $\hat{\delta} = 0.23$ for the degree of path dependence. In addition, $\hat{\chi}$ is estimated as 0.87 (standard error = 0.31), which maps to a long-run effect of 1.14—hence, 1 percent more aging will be associated with 1.14 percent shifts toward automation technologies. These

estimates also imply an elasticity of substitution between factors of $\hat{\sigma} = 1.69$, which ensures that $\hat{\delta}\hat{\sigma} = 0.40 < 1$.

These parameters, together with equation (21), yield a lower bound distortion of $n^{SP}/n^{EQ} = 0.82$, as shown in panel C at the bottom of the table. This is a sizable difference between the equilibrium and socially optimal direction of technology—a socially optimal technology ratio that is 18 percent lower than the equilibrium—despite the fact that the pecuniary externality in the automation case appears small. This magnitude is partly explained by the nontrivial value of $\hat{\delta}\hat{\sigma} = 0.40$, which amplifies the impact of distortions. Nevertheless, the welfare loss from equilibrium distortions is modest, about 1 percent in consumption-equivalent terms. Panel D shows that using an even smaller estimate of quasi-rents from employment (7.5 percent instead of 15 percent) gives correspondingly smaller numbers for the technology distortion ($n^{SP}/n^{EQ} = 0.91$) and welfare losses (0.2 percent).

Column 2 of Table 2 considers one variation on the automation numbers by using ten-year rather than five-year intervals. The results are broadly similar: $\hat{\rho} = 0.76$ (standard error = 0.12), $\hat{\chi} = 1.16$ (standard error = 0.38), and a long-run effect of 1.52. These imply $\hat{\sigma} = 1.85$ and $\hat{\delta}\hat{\sigma} = 0.44$, which together yield slightly larger technology distortions and welfare costs: $n^{SP}/n^{EQ} = 0.79$ and 1 percent in panel C. Panel D numbers are correspondingly smaller.

Table B1 in online Appendix B presents a number of robustness checks and additional results. In particular, in columns 3–8, I show that similar results hold when instead of $\ln x$, I use $\ln(1+x)$ and include observations with zeros; when I use the inverse hyperbolic sine, $a \sinh$ (a transformation that allows for zeros and approximately yields logarithmic form for nonzero observations); and for the OECD sample. The implied technology distortion n^{SP}/n^{EQ} remains comparable to those in columns 1 and 2, ranging from 0.56 to 0.76 in panel C. The exceptions are the inverse hyperbolic sine model and the specification that focuses on just OECD countries, in both cases at the five-year horizon (columns 5 and 7). In these instances, the estimates for δ are higher, and consequently technology distortions are more pronounced ($n^{SP}/n^{EQ} = 0.40$ and 0.34).

²² When using five-year (ten-year) changes, the anticipated aging variable is for the next 15 (20) years. Acemoglu and Restrepo (2022) also show that instrumental variable estimates exploiting fertility changes from several decades before give very similar results to these ordinary least squares (OLS) estimates. Here, I focus on OLS models.

TABLE 2—ESTIMATES AND IMPLIED PARAMETERS

Application	Automation		Health		Energy	
	5-year (1)	10-year (2)	5-year (3)	10-year (4)	5-year (5)	10-year (6)
<i>Panel A. Parameters estimated from regressions</i>						
Initial relative stock: $\hat{\rho}$	0.77 (0.14)	0.76 (0.12)	0.93 (0.03)	1.11 (0.03)	0.81 (0.03)	0.86 (0.04)
Initial shifter: $\hat{\chi}$	0.87 (0.31)	1.16 (0.38)	0.10 (0.01)	0.14 (0.01)	−1.52 (0.29)	−1.06 (0.66)
Changes in shifter: $\hat{\lambda}$	1.18 (0.43)	1.81 (0.63)	−0.004 (0.02)	0.001 (0.02)	−0.45 (0.20)	1.12 (0.82)
Observations	232	125	55,699	37,389	13,648	6,824
<i>Panel B. Implied parameters</i>						
Long-run effects	1.14	1.52	0.11	0.14	−1.89	−1.23
$\hat{\delta}$	0.23	0.24	0.07	0.00	0.19	0.14
$\hat{\sigma}$	1.69	1.85	1.10	1.14	2.73	2.53
$\hat{\varepsilon}$	4.09	4.82	1.18	1.26	7.27	6.56
$\hat{\delta}\hat{\sigma}$	0.40	0.44	0.08	0.00	0.53	0.36
<i>Panel C. Equilibrium and welfare comparison</i>						
n^{SP}/n^{EQ}	0.82	0.79	0.43	0.45	0.44	0.57
$U^{SP} - U^{EQ}$	0.01	0.01	0.06	0.05	0.03	0.02
<i>Panel D. Equilibrium and welfare comparison (alternatives)</i>						
n^{SP}/n^{EQ}	0.91	0.89	0.58	0.59	0.00	0.00
$U^{SP} - U^{EQ}$	0.002	0.002	0.18	0.17	13.74	8.94

Notes: This table presents regression estimates (panel A), implied parameter values (panel B), and implied distortions and welfare results (panels C and D) for the three applications. In all cases, regressions are estimated with OLS, and heteroscedasticity-robust standard errors are presented in parentheses. Standard errors are clustered at the country level in columns 1–4 and at the firm level in columns 5 and 6. Odd-numbered columns are for five-year changes, and even-numbered columns are for ten-year changes. Columns 1 and 2 are for the automation application and are at the country-time-period level and present regressions weighted by manufacturing employment in 1990. The dependent variable is the number of newly granted patents for automation technologies relative to other utility patents divided by the stock of patents related to automation relative to the stock of other utility patents (in logs). Shifters are the level and change in the ratio of workers above the age of 56 to workers between 21 and 55. Column 1 uses expected 20-year change, and column 2 uses expected 15-year change (in logs). Both columns include region dummies and the 1990 values of log GDP per capita, log of population, average years of schooling, and the ratio of workers above age 56 to workers aged 21 in 1990 interacted with period dummies. Columns 3 and 4 are for the health-care application, and observations are at the country-disease-period level. The dependent variable is relative number of new medical articles for each disease divided by relative stock of medical articles for that disease (in logs). Shifters are the log of the burden of disease for the relevant country-year-disease cell. Both columns include country, disease, and period fixed effects. Columns 5 and 6 are for the energy application, and observations are at the firm-period level. The dependent variable is relative number of newly granted patents for dirty technologies relative to newly granted patents for clean technologies (with the $\log(1+x)$ transformation). Shifters are firm-level fuel prices adjusted (based on firm-level fuel consumption) inclusive of taxes. Both columns include firm and period fixed effects as well as the values of government R&D subsidies for clean innovation, regulations over emissions, and the relevant country's GDP per capita for that period (as in Aghion et al. 2016). In columns 1 and 2, panel C uses 15 percent quasi-rents for workers, and panel D uses 7.5 percent quasi-rents. In columns 3 and 4, panel C focuses on markup differences, and panel D uses the externality estimate computed from the shortfall of QALYs from curative relative to preventative technologies. In columns 5 and 6, panels C and D use externality numbers based on Rennert et al.'s (2022) estimate of the social cost of CO₂ for the United States and the world, respectively.

and welfare losses are also larger. Finally, columns 9 and 10 of this table report estimates of equation (23) from the recent paper by Dechezleprêtre et al. (2021), who study the effects of skill premia on automation

technologies at the firm level. Using five-yearly observations across about 1,150 firms that have at least four automation patents, these columns show similar estimates of the degree of path dependence and the elasticity of substitution σ

to the baseline estimates in columns 1 and 2 of Table 2 (column 9 includes firm fixed effects and industry-by-time fixed effects, while column 10 additionally includes country-by-time fixed effects). As a result, we obtain broadly comparable technology distortions using estimates from this firm-level dataset: $n^{SP}/n^{EQ} = 0.47$ and 0.61 in the two columns, with welfare losses of 3 percent and 2 percent, respectively.

D. Estimates: Health

Because detailed data classified into preventative and curative health innovations are not available, for the regression analysis I use data on medical research and disease burdens from Acemoglu et al. (2023). Columns 3 and 4 report estimates from equation (24) using these data. The left-hand-side variable is the flow of medical articles for a disease in a country during a particular time period (relative to the stock of medical articles relevant for this observation), and the forcing variable is the disease burden for that disease, country, and time. Disease burdens are computed as declines in the number of disability-adjusted life years caused by each disease in a country and time period in our sample.²³ All regressions in this case are unweighted and control for disease, country, and time fixed effects.

Column 3 focuses on five-year periods, while column 4 looks at ten-year observations. In both columns, the sample covers the years 1990–2019 and 279 diseases and comes from 193 countries. In column 3, we have a total of 55,699 observations, while there are 37,389 observations in column 4.

The estimates in the two columns are similar. In column 3, $\hat{\rho}$ is 0.93 (standard error = 0.03), which implies a path dependence parameter of $\hat{\delta} = 0.07$. The estimate of $\hat{\chi} = 0.10$ combined with these numbers yields a long-run effect of 0.11. Hence, a 1 percent increase in the burden of a specific disease in a country leads to a 0.11 percent increase in the medical research directed to that disease. The implied elasticity of substitution is $\hat{\sigma} = 1.10$, which again puts us comfortably in the region where

$\hat{\delta}\hat{\sigma} = 0.08 < 1$. In column 4, ρ is estimated to be a little more than 1 (1.11), which implies no path dependence, and thus I set $\delta = 0$. Other estimates remain similar—in particular, a long-run effect of 0.14 and $\sigma = 1.14$.

Panel C focuses on markup differences between preventative and curative categories, given in Table 1 (and ignores differences in externalities). Technology distortions are similar in the two columns: $n^{SP}/n^{EQ} = 0.43$ in column 3 and 0.45 in column 4, meaning that the technology ratio is about 45 percent biased in favor of curative technologies in the decentralized equilibrium. The resulting welfare effects are sizable—around 5–6 percent (which should be interpreted as a fraction of health-care consumption).

Panel D looks at the implications of the τ_2 estimate from the shortfall of QALY gains from curative technologies relative to preventative technologies (now ignoring markup differences). This alternative way of conceptualizing misaligned innovation incentives in health care leads to even larger technology distortions: n^{SP}/n^{EQ} is around 0.6, and welfare losses from the equilibrium direction of technology are correspondingly bigger (17–18 percent).

In Table B2 in online Appendix B, I show that the estimates reported in columns 3 and 4 of Table 2, and thus the implied technology distortions and welfare effects, are quite robust. Similar results are obtained when instead of $\ln x$, I use $\ln(1+x)$ and keep observations with zeros; when I use the inverse hyperbolic sine ($a \sinh$) transformation; when the country fixed effects are omitted; when we include country-times-year and disease-times-year fixed effects; and when we focus only in variation in the United States. The implied values for n^{SP}/n^{EQ} in panel C are mostly around 0.4, and the welfare effects are also comparable to those in Table 2, except in specifications using $\ln(1+x)$ and $a \sinh$ with five-year observations and in the two specifications that do not include country fixed effects, where technology distortions are larger, ranging between 0.11 and 0.17, and the welfare effects are correspondingly more substantial.

E. Estimates: Energy

In the context of energy, I follow the conceptual structure in Acemoglu et al. (2012) that distinguishes dirty (coal, gas, and oil) technologies and clean (renewables and nuclear)

²³ These calculations are based on data from the Global Burden of Disease project, which is a collaboration between the World Bank and the Institute for Health Metrics and Evaluation. See Acemoglu et al. (2023) for details.

technologies. Columns 5 and 6 of Table 2 use data from Aghion et al. (2016) and report firm-level regressions of the flow of patents of clean technologies relative to dirty technologies in the automobile sector, once again normalized with their respective stocks. In these data, there are many observations with zero stocks, and I follow Aghion et al. (2016) and include these observations by using $\ln(1+x)$. This gives a dataset of 3,412 firms across 58 countries for the years 1986–2005 and 13,684 and 6,824 observations in the two columns.

The estimates are fairly similar between columns 5 and 6. For example, ρ is estimated as 0.81 (standard error = 0.03) in column 5 and 0.86 (standard error = 0.04) in column 6. Long-run effects are comparable as well: -1.89 in column 5 and -1.23 in column 6 (these are the effects of higher gasoline prices, leading to lower clean technology patents, hence the negative sign). The estimated values of σ are also similar across the two columns: 2.73 and 2.53. As a result, in both columns, we have $\hat{\delta}\hat{\sigma} < 1$.

Using our baseline estimate of $\tau_2 = 0.13$ in panel C of Table 1 based on social cost of carbon for the United States, the technology distortion is found to be $n^{SP}/n^{EQ} = 0.44$ in column 5 and a little smaller, $n^{SP}/n^{EQ} = 0.57$, in column 6. These are again sizable distortions with welfare losses of about 2–3 percent.

Instead, with global damages, we have $\tau_2 = 0.94$, and because this externality is close to 1, equation (20) implies that the social planner would like to essentially shut down fossil fuel technologies (i.e., $n^{SP}/n^{EQ} \approx 0$), as indicated in panel D.

In Table B3 in online Appendix B, I report various robustness checks. The general pattern is broadly comparable to that shown in columns 5 and 6. Long-run effects and elasticity estimates are quite similar, including in specifications that add spillovers from the stock of innovation of other firms in the same country, as in Aghion et al. (2016). The extent of technology distortions, n^{SP}/n^{EQ} , remains fairly stable and ranges between 0.37 and 0.74 across all specifications in panel C.

Overall, in all three of the applications considered here, I find suggestive evidence that distortions in the direction of technology can be sizable and generate nontrivial welfare consequences. These results should be interpreted with ample caution since both the estimates of

the underlying parameters and even more so the estimates of externalities and markups are subject to considerable uncertainty. They are presented in the spirit of suggestive evidence to stimulate more work in this area.

V. Concluding Remarks

Technological change is vital for continued economic prosperity and can help tackle many of the epochal challenges facing humanity, such as climate change, pandemics, and global poverty. Because of its society-wide benefits, corporations and individuals tend to underinvest in innovation, and this underinvestment provides a central justification for government support for science, academia, and corporate R&D. But will the “market process”—working through profit incentives, competition, and reputational concerns of researchers—get the direction of innovation right?

Typically, there are many alternative technologies and paradigms even within a narrow field. In health care, innovation can be directed toward curative technologies and pharmaceuticals, or it can prioritize preventative technologies. In energy and transport, innovation can be directed toward clean or dirty alternatives. In most industries, researchers and corporations decide how much to invest to automate production processes versus how much to prioritize increasing worker marginal productivity by providing better tools, new labor-intensive tasks, and new learning opportunities to employees. In agriculture, novel crop varieties can target pests and pathogens that are pervasive in some countries ahead of others.

In this paper, I have suggested that there may be systemic reasons for the direction of innovation to be distorted. Using a simple framework, I highlighted the factors impacting the direction of technology and illustrated how economic or social externalities (such as carbon emissions) and markup differences between technologies can lead to a misaligned direction of innovation. Innovation distortions tend to reduce or even reverse welfare gains from technological progress (e.g., when research effort focuses on socially costly technologies) and can even slow down economic growth (e.g., because of markup differences).

There are three distinct objections one could raise to the approach in this paper. First, even if the market does not get the direction of

innovation completely right, governments and bureaucrats could be worse at it. This objection is valid and is the reason why much of my discussion focused on systemic sources of distortions that can be determined without superior technical knowledge on the part of bureaucrats or some impressive ability to “pick winners.” If there are markup differences across the products generated by different technologies or quantifiable externalities—as I have proposed—then the extent of distortions can be determined and agreed upon.

Second, one may argue that distortions resulting from the direction of technology are secondary relative to underinvestment in overall innovation and/or they are small relative to other costs that government intervention in the innovation process would generate. This is also a valid concern, but ultimately the extent of these distortions is a quantitative question. For this reason, I provided evidence from three distinct domains on distortions in the direction of technology.

Third, attempts to deal with distortions in the direction of innovation could lead to new and challenging political economy questions. I return to this important question at the end of these remarks.

In light of these caveats, the current paper should be seen as a first step in a more detailed investigation of possible distortions in the direction of technological change and potential remedies. This is the reason why the theoretical framework is chosen to be as simple as possible and the quantitative evaluation is purely suggestive. Several interesting questions are open for future study within this framework, and I list some of them here.

It would be instructive to model and empirically investigate the extent to which other social factors can also create distortions in the direction of scientific and corporate research. One possibility is researchers following each other’s leads and becoming influenced by each other’s visions to such an extent that it makes them overinvest in some paradigms. I have suggested in past work (Acemoglu and Restrepo 2020b; Acemoglu and Johnson 2023) that this may be a concern within the field of AI, pushing researchers to prioritize automation and mass-scale data collection. What the theoretical microfoundations of such effects are, whether this type of bias is indeed present in practice, and whether

governmental or societal intervention may be possible in this case are interesting questions for future research.

The theoretical analysis in the paper ignored the interplay between Pigovian taxes and policy aimed at redirecting technological change. A critical question from both a theoretical and an applied point of view is to what extent these different classes of policies are complements or substitutes.

Much industrial policy became mired in corruption and political problems in the past, and one may be worried that any government intervention aimed at influencing the direction of technological change would be similarly hampered by political economy challenges. This is particularly true since history is full of examples of special interest groups attempting to block technological change to protect their rents or privileges (e.g., Acemoglu and Robinson 2012). On the other hand, the endogeneity of the direction of innovation opens up new political economy avenues, and studying them is a fruitful area for future inquiry (see Acemoglu and Johnson 2023).

In this context, another research area is to model the market structure of the relevant industries in greater detail so that the pro- or anticompetitive effects of policies aimed at redirecting technological change can be evaluated—for example, can firms and researchers be encouraged to invest in socially more beneficial technologies without reducing the extent of competition in the economy?

Finally, the empirical part of the current paper was a first attempt, and more systematic work on measuring distortions in the direction of innovation is a critical area for future research.

REFERENCES

- Acemoglu, Daron. 1998. “Why Do New Technologies Complement Skills? Directed Technical Change and Wage Inequality.” *Quarterly Journal of Economics* 113 (4): 1055–90.
- Acemoglu, Daron. 2002. “Directed Technical Change.” *Review of Economic Studies* 69 (4): 781–809.
- Acemoglu, Daron. 2003a. “Labor- and Capital-Augmenting Technical Change.” *Journal of the European Economic Association* 1 (1): 1–37.

- Acemoglu, Daron.** 2003b. "Patterns of Skill Premia." *Review of Economic Studies* 70 (2): 199–230.
- Acemoglu, Daron.** 2007. "Equilibrium Bias of Technology." *Econometrica* 75 (5): 1371–409.
- Acemoglu, Daron.** 2010. "When Does Labor Scarcity Encourage Innovation?" *Journal of Political Economy* 118 (6): 1037–78.
- Acemoglu, Daron.** 2012. "Diversity and Technological Progress." In *The Rate and Direction of Inventive Activity Revisited*, edited by Josh Lerner and Scott Stern, 319–56. Chicago: University of Chicago Press.
- Acemoglu, Daron.** 2023. "Replication data for: Distorted Innovation: Getting the Direction of Technology Wrong." American Economic Association [publisher], Inter-university Consortium for Political and Social Research [distributor]. <https://doi.org/10.38886/E184829V1>.
- Acemoglu, Daron, Philippe Aghion, Leonardo Bursztyn, and David Hémous.** 2012. "The Environment and Directed Technical Change." *American Economic Review* 102 (1): 131–66.
- Acemoglu, Daron, Nicholas Ajzeman, Cevat Giray Aksoy, Martin Fiszbein, and Carlos Molina.** 2021. "(Successful) Democracies Read Their Own Support." NBER Working Paper 29167.
- Acemoglu, Daron, Ufuk Akcigit, Douglas Hanley, and William Kerr.** 2016. "Transition to Clean Technology." *Journal of Political Economy* 124 (1): 52–140.
- Acemoglu, Daron, Ufuk Akcigit, and William Kerr.** 2016. "Innovation Network." *Proceedings of the National Academy of Sciences* 113 (41): 11483–88.
- Acemoglu, Daron, and David Autor.** 2011. "Skills, Tasks and Technologies: Implications for Employment and Earnings." In *Handbook of Labor Economics*, Vol. 4, edited by David Card and Orley Ashenfelter. Amsterdam: Elsevier.
- Acemoglu, Daron, David Cutler, Amy Finkelstein, and Joshua Linn.** 2006. "Did Medicare Induce Pharmaceutical Innovation?" *American Economic Review* 96 (2): 103–07.
- Acemoglu, Daron, David Hémous, Lint Barrage, and Philippe Aghion.** 2019. "Climate Change, Directed Innovation, and Energy Transition: The Long-Run Consequences of the Shale Gas Revolution." Unpublished.
- Acemoglu, Daron, and Simon Johnson.** 2023. *Power and Progress: Our Thousand-Year Struggle over Technology and Prosperity*. New York: Public Affairs.
- Acemoglu, Daron, Michael I. Jordan, and Glen Weyl.** 2021. "The Turing Test Is Bad for Business." *Wired*, November 8. <https://www.wired.com/story/artificial-intelligence-turing-test-economics-business/>.
- Acemoglu, Daron, and Todd Lensman.** 2023. "Technology Paradigms, Lock-in, and Economic Growth." Unpublished.
- Acemoglu, Daron, and Joshua Linn.** 2004. "Market Size in Innovation: Theory and Evidence from the Pharmaceutical Industry." *Quarterly Journal of Economics* 119 (3): 1049–90.
- Acemoglu, Daron, Andrea Manera, and Pascual Restrepo.** 2020. "Does the US Tax Code Favor Automation?" *Brookings Papers on Economic Activity* 50 (1): 231–85.
- Acemoglu, Daron, Jacob Moscona, Karthik A. Sastry, and Heidi Williams.** 2023. "The Rich-World Bias of Medical Research." Unpublished.
- Acemoglu, Daron, and Pascual Restrepo.** 2018. "The Race between Man and Machine: Implications of Technology for Growth, Factor Shares, and Employment." *American Economic Review* 108 (6): 1488–542.
- Acemoglu, Daron, and Pascual Restrepo.** 2020a. "Robots and Jobs: Evidence from US Labor Markets." *Journal of Political Economy* 128 (6): 2188–244.
- Acemoglu, Daron, and Pascual Restrepo.** 2020b. "The Wrong Kind of AI? Artificial Intelligence and the Future of Labour Demand." *Cambridge Journal of Regions, Economy and Society* 13 (1): 25–35.
- Acemoglu, Daron, and Pascual Restrepo.** 2022. "Demographics and Automation." *Review of Economic Studies* 89 (1): 1–44.
- Acemoglu, Daron, and James A. Robinson.** 2012. *Why Nations Fail: The Origins of Power, Prosperity and Poverty*. New York: Crown Business.
- Acemoglu, Daron, David Y. Yang, and Jie Zhou.** 2023. "Power and the Direction of Research: Evidence from China's Academia." Unpublished.
- Acemoglu, Daron, and Fabrizio Zilibotti.** 2001. "Productivity Differences." *Quarterly Journal of Economics* 116 (2): 563–606.
- Aghion, Philippe, Antoine Dechezleprêtre, David Hémous, Ralf Martin, and John Van Reenen.** 2016. "Carbon Taxes, Path Dependency, and Directed Technical Change: Evidence from the

- Auto Industry.” *Journal of Political Economy* 124 (1): 1–51.
- Agnell, Marcia.** 2004. *The Truth About Drug Companies: How They Deceive Us and What to Do about It*. New York: Random House.
- Ahmad, Syed.** 1966. “On the Theory of Induced Invention.” *Economic Journal* 76 (302): 344–57.
- Akcigit, Ufuk, Douglas Hanley, and Nicolas Serano-Velarde.** 2021. “Back to Basics: Basic Research Spillovers, Innovation Policy, and Growth.” *Review of Economic Studies* 88 (1): 1–43.
- Allen, Robert C.** 2009. *The British Industrial Revolution in Global Perspective*. Cambridge: Cambridge University Press.
- Anderson, Gerard F., Peter Hussey, and Varduhi Petrosyan.** 2019. “It’s Still the Prices, Stupid: Why the US Spends So Much on Health Care, and a Tribute to Uwe Reinhardt.” *Health Affairs* 38 (1): 87–95.
- Arthur, W. Brian.** 1989. “Competing Technologies, Increasing Returns, and Lock-In by Historical Events.” *Economic Journal* 99 (394): 116–31.
- Autor, David H., David Dorn, and Gordon Hanson.** 2019. “When Work Disappears: How Adverse Labor Market Shocks Affect Fertility, Marriage, and Children’s Living Circumstances.” *American Economic Review: Insights* 1 (2): 161–78.
- Autor, David H., Frank Levy, and Richard J. Murnane.** 2003. “The Skill Content of Recent Technological Change: An Empirical Exploration.” *Quarterly Journal of Economics* 118 (4): 1279–333.
- Azoulay, Pierre, Joshua S. Graff Zivin, Danielle Li, and Bhaven N. Sampat.** 2019. “Public R&D Investments and Private-Sector Patenting: Evidence from NIH Funding Rules.” *Review of Economic Studies* 86 (1): 117–52.
- Bloom, Nicholas, Mark Schankerman, and John Van Reenen.** 2013. “Identifying Technology Spillovers and Product Market Rivalry.” *Econometrica* 81 (4): 1347–93.
- Bovenberg, A. Lans, and Sjak Smulders.** 1995. “Environmental Quality and Pollution-Augmenting Technological Change in a Two-Sector Endogenous Growth Model.” *Journal of Public Economics* 57 (3): 369–91.
- Branstetter, Lee G., Guangwei Li, and Mengjia Ren.** 2022. “Picking Winners? Government Subsidies and Firm Productivity in China.” NBER Working Paper 30699.
- Brynjolfsson, Erik.** 2022. “The Turing Trap: The Promise & Peril of Human-Like Artificial Intelligence.” *Dædalus* 151 (2): 272–87.
- Budish, Eric, Benjamin N. Roin, and Heidi Williams.** 2015. “Do Firms Underinvest in Long-Term Research? Evidence from Cancer Clinical Trials.” *American Economic Review* 105 (7): 2044–85.
- Case, Anne, and Angus Deaton.** 2020. *Deaths of Despair and the Future of Capitalism*. Princeton, NJ: Princeton University Press.
- Clemens, Michael A., Ethan G. Lewis, and Hannah M. Postel.** 2018. “Immigration Restrictions as Active Labor Market Policy: Evidence from the Mexican Bracero Exclusion.” *American Economic Review* 108 (6): 1468–87.
- Cockburn, Iain, and Zvi Griliches.** 1988. “Industry Effects and Appropriability Measures in the Stock Market’s Valuation of R&D and Patents.” *American Economic Review* 78 (2): 419–23.
- Coe, David T., and Elhanan Helpman.** 1995. “International R&D Spillovers.” *European Economic Review* 39 (5): 859–87.
- Costinot, Arnaud, Dave Donaldson, Margaret Kyle, and Heidi Williams.** 2019. “The More We Die, The More We Sell? A Simple Test of the Home-Market Effect.” *Quarterly Journal of Economics* 134 (2): 843–94.
- Dechezleprêtre, Antoine, David Hémous, Morten Olsen, and Carlo Zanella.** 2021. “Induced Automation: Evidence from Firm-Level Patent Data.” Unpublished.
- Dechezleprêtre, Antoine, Ralf Martin, and Myra Mohnen.** 2014. “Knowledge Spillovers from Clean and Dirty Technologies.” Unpublished.
- De Loecker, Jan, Jan Eeckhout, and Gabriel Unger.** 2020. “The Rise of Market Power and the Macroeconomic Implications.” *Quarterly Journal of Economics* 135 (2): 561–644.
- Di Maria, Corrado, and Simone Valente.** 2008. “Hicks Meets Hotelling: The Direction of Technical Change in Capital-Resource Economies.” *Environment and Development Economics* 13 (6): 691–717.
- Diao, Xinshen, Mia Ellis, Margaret S. McMillan, and Dani Rodrik.** 2021. “Africa’s Manufacturing Puzzle: Evidence from Ethiopian and Tanzanian Firms.” NBER Working Paper 28344.
- Dosi, Giovanni.** 1982. “Technological Paradigms and Technological Trajectories: A Suggested Interpretation of the Determinants and

- Directions of Technical Change.” *Research Policy* 11 (3): 147–62.
- Drandakis, E.M., and E.S. Phelps.** 1966. “A Model of Induced Invention, Growth and Distribution.” *Economic Journal* 76 (304): 823–40.
- Finkelstein, Amy.** 2004. “Static and Dynamic Effects of Health Policy: Evidence from the Vaccine Industry.” *Quarterly Journal of Economics* 119 (2): 527–64.
- Gancia, Gino, and Fabrizio Zilibotti.** 2005. “Horizontal Innovation in the Theory of Growth and Development.” In *Handbook of Economic Growth*, Vol. 1A, edited by Philippe Aghion and Steven N. Durlauf, 111–70. Amsterdam: Elsevier.
- Goulder, Lawrence H., and Stephen H. Schneider.** 1999. “Induced Technological Change and the Attractiveness of CO2 Abatement Policies.” *Resource and Energy Economics* 21 (3–4): 211–53.
- Grimaud, André, and Luc Rouge.** 2008. “Environment, Directed Technical Change and Economic Policy.” *Environmental and Resource Economics* 41: 439–63.
- Grossman, Gene M., and Elhanan Helpman.** 1993. *Innovation and Growth in the Global Economy*. Cambridge, MA: MIT Press.
- Gruber, Jonathan, and Simon Johnson.** 2019. *Jump-Starting America: How Breakthrough Science Can Revive Economic Growth and the American Dream*. New York: PublicAffairs.
- Habakkuk, H. J.** 1962. *American and British Technology in the Nineteenth Century: Search for Labor Saving Inventions*. Cambridge, UK: Cambridge University Press.
- Hager, Thomas.** 2009. *The Alchemy of Air: A Jewish Genius, A Doomed Tycoon, and the Scientific Discovery That Fed the World but Fueled the Rise of Hitler*. New York: Broadway Books.
- Hanlon, W. Walker.** 2015. “Necessity is the Mother of Invention: Input Supplies and Directed Technical Change.” *Econometrica* 83 (1): 67–100.
- Hayami, Yujiro, and V.W. Ruttan.** 1970. “Factor Prices and Technical Change in Agricultural Development: The United States and Japan, 1880–1960.” *Journal of Political Economy* 78 (5): 1115–41.
- Hémous, David.** 2016. “The Dynamic Impact of Unilateral Environmental Policies.” *Journal of International Economics* 103: 80–95.
- Hémous, David, and Morten Olsen.** 2022. “The Rise of the Machines: Automation, Horizontal Innovation, and Income Inequality.” *American Economic Journal: Macroeconomics* 14 (1): 179–223.
- Hicks, John.** 1932. *The Theory of Wages*. London: Macmillan.
- Hopenhayn, Hugo, and Francesco Squintani.** 2021. “On the Direction of Innovation.” *Journal of Political Economy* 129 (7): 1991–2022.
- Hornbeck, Richard.** 2012. “The Enduring Impact of the American Dust Bowl: Short- and Long-Run Adjustments to Environmental Catastrophe.” *American Economic Review* 102 (4): 1477–507.
- Howard, David H., Peter B. Bach, Ernst R. Berndt, and Rena M. Conti.** 2015. “Pricing in the Market for Anticancer Drugs.” *Journal of Economic Perspectives* 29 (1): 139–62.
- Howell, Sabrina T.** 2017. “Financing Innovation: Evidence from R&D Grants.” *American Economic Review* 107 (4): 1136–64.
- Howell, Sabrina T., Jason Rathje, John Van Reenen, and Jun Wong.** 2021. “Opening Up Military Innovation: Causal Effects of ‘Bottom-up’ Reforms to U.S. Defense Research.” Unpublished.
- Hufbauer, Gary Clyde, and Euijin Jung.** 2021. *Scoring 50 Years of US Industrial Policy, 1970–2020*. Washington, DC: Peterson Institute for International Economics.
- Jones, Charles I.** 2005. “The Shape of Production Functions and the Direction of Technical Change.” *Quarterly Journal of Economics* 120 (2): 517–49.
- Jones, Charles I., and John C. Williams.** 1998. “Measuring the Social Return to R&D.” *Quarterly Journal of Economics* 113 (4): 1119–35.
- Kennedy, Charles.** 1964. “Induced Bias in Innovation and the Theory of Distribution.” *Economic Journal* 74 (295): 541–47.
- Kenkel, Donald S.** 2000. “Prevention.” In *Handbook of Health Economics*, Vol. 1, edited by Anthony J. Culyer and Joseph P. Newhouse, 1675–720. Amsterdam: Elsevier.
- Kiley, M.T.** 1999. “The Supply of Skilled Labour and Skill-Biased Technological Progress.” *Economic Journal* 109 (458): 708–24.
- Koyama, Mark, and Jared Rubin.** 2022. *How the World Became Rich: The Historical Origins of Economic Growth*. Cambridge, UK: Polity Press.
- Kremer, Michael, and Christopher M. Snyder.** 2015. “Preventives versus Treatments.” *Quarterly Journal of Economics* 130 (3): 1167–239.

- Lane, Nathan.** 2021. "Manufacturing Revolutions: Industrial Policy and Industrialization in South Korea." Unpublished.
- Martin, Ralf, and Dennis Verhoeven.** 2022. "Knowledge Spillovers from Clean and Emerging Technologies in the UK." Paper presented at International Association for Research in Income and Wealth 37th General Conference, Luxembourg, August 22–26, 2022. <https://iariw.org/wp-content/uploads/2022/08/Martin-Verhoeven-IARIW-2022.pdf>.
- Mazzucato, Mariana.** 2015. *The Entrepreneurial State: Debunking Public vs. Private Sector Myths*. New York: Anthem Press.
- Mitrunen, Matti.** 2019. "War Reparations, Structural Change, and Intergenerational Mobility." Unpublished.
- Mokyr, Joel.** 1992. *The Levers of Riches: Technological Creativity and Economic Progress*. New York: Oxford University Press.
- Mokyr, Joel.** 2004. *The Gifts of Athena*. Princeton, NJ: Princeton University Press.
- Moretti, Enrico, Claudia Steinwender, and John Van Reenen.** 2019. "The Intellectual Spoils of War? Defense R&D, Productivity and International Spillovers." NBER Working Paper 26483.
- Moscona, Jacob.** 2021. "Environmental Catastrophe and the Direction of Invention: Evidence from the American Dust Bowl." Unpublished.
- Moscona, Jacob, and Karthik A. Sastry.** Forthcoming. "Does Directed Innovation Mitigate Climate Damage? Evidence from U.S. Agriculture." *Quarterly Journal of Economics*.
- Moscona, Jacob, and Karthik A. Sastry.** 2023. "Inappropriate Technology: Evidence from Global Agriculture." Unpublished.
- Newell, Richard G., Adam B. Jaffe, and Robert N. Stavins.** 1999. "The Induced Innovation Hypothesis and Energy-Saving Technological Change." *Quarterly Journal of Economics* 114 (3): 941–75.
- Newhouse, Joseph P.** 2021. "An Ounce of Prevention." *Journal of Economic Perspectives* 35 (2): 101–18.
- Pack, Howard, and Kamal Saggi.** 2006. "Is There a Case for Industrial Policy? A Critical Survey." *World Bank Research Observer* 21 (2): 267–97.
- Popp, David.** 2002. "Induced Innovation and Energy Prices." *American Economic Review* 92 (1): 160–80.
- Rennert, Kevin, Frank Errickson, Brian C. Prest, Lisa Rennels, Richard G. Newell, William Pizer, Cora Kingdon, et al.** 2022. "Comprehensive Evidence Implies a Higher Social Cost of CO₂." *Nature* 610 (7933): 687–92.
- Rodrik, Dani.** 2014. "Green Industrial Policy." *Oxford Review of Economic Policy* 30 (3): 469–91.
- Ridley, Matt.** 2020. *How Innovation Works: And Why It Flourishes in Freedom*. New York: Harper.
- Samuelson, Paul A.** 1965. "A Theory of Induced Innovation along Kennedy-Weisäcker Lines." *Review of Economics and Statistics* 47 (4): 343–56.
- Thoenig, Mathias, and Thierry Verdier.** 2003. "A Theory of Defensive Skill-Biased Innovation and Globalization." *American Economic Review* 93 (3): 709–28.
- Wilson, William Julius.** 1996. *When Work Disappears: The World of the New Urban Poor*. New York: Alfred A. Knopf.
- Zeira, Joseph.** 1998. "Workers, Machines, and Economic Growth." *Quarterly Journal of Economics* 113 (4): 1091–117.
- Zuckerman, Gregory.** 2021. *A Shot to Save the World: The Inside Story of the Life-or-Death Race for a COVID-19*. Brentford: Portfolio.

This article has been cited by:

1. Keith O. Fuglie, Thomas W. Hertel, David B. Lobell, Nelson B. Villoria. 2024. Agricultural Productivity and Climate Mitigation. *Annual Review of Resource Economics* **16**:1, 21-40. [[Crossref](#)]
2. Dibyendu Maiti, Chiranjib Neogi. 2024. Endogeneity-corrected stochastic frontier with market imperfections. *Empirical Economics* **67**:3, 1149-1183. [[Crossref](#)]
3. Esmée S. R. Dijk, José L. Moraga-González, Evgenia Motchenkova. 2024. Start-up acquisitions, strategic R&D, and the entrant's and incumbent's direction of innovation. *Journal of Economics & Management Strategy* **36**. . [[Crossref](#)]
4. Bart van Ark, Klaas de Vries, Dirk Pilat. 2023. Are pro-productivity policies fit for purpose?. *The Manchester School* **21**. . [[Crossref](#)]